


Table of Contents

1. [Data Lineage • Overview • Palantir](#)
2. [Data Lineage • Overview • Palantir](#)
3. [Data Lineage • Navigation • Palantir](#)
4. [Data Lineage • Branching data lineage • Palantir](#)
5. [Data Lineage • FAQ • Palantir](#)
6. [Data Lineage • Graphs • Explore data lineage • Palantir](#)
7. [Data Lineage • Graphs • Explore artifacts and ontology entities • Palantir](#)
8. [Data Lineage • Graphs • Save and share a graph • Palantir](#)
9. [Data Lineage • Graphs • Node coloring • Palantir](#)
10. [Data Lineage • Graphs • Graph elements reference • Palantir](#)
11. [Data Lineage • Understand and manage datasets • View dataset preview and logic • Palantir](#)
12. [Data Lineage • Understand and manage datasets • View build timeline • Palantir](#)
13. [Data Lineage • Understand and manage datasets • Understand out-of-date datasets • Palantir](#)
14. [Data Lineage • Understand and manage datasets • Find datasets with a given column • Palantir](#)
15. [Data Lineage • Understand and manage datasets • Build datasets • Palantir](#)
16. [Data Lineage • Understand and manage datasets • Manage schedules • Palantir](#)
17. [Data Lineage • Understand and manage datasets • Roll back a pipeline • Palantir](#)
18. [Data Lineage • Understand and manage datasets • Roll back a dataset • Palantir](#)
19. [Data Lineage • Understand permissions • Check resource permissions • Palantir](#)
20. [Data Lineage • Understand permissions • See the impact of marking changes • Palantir](#)

Data connectivity & integration / Data Lineage / Overview

Data Lineage

Data Lineage is an interactive tool that facilitates a holistic view of how data flows through the Foundry platform.



With Data Lineage, you can:

- Easily find and discover datasets
 - Search to find datasets using Project, table, and column names
 - Click through Foundry Projects to browse data
- Explore [pipelines](#) through a powerful interface
 - Expand or hide ancestors and descendants of datasets

API Reference ↗ s of a group of tables at once

Visualize your pipeline through coloring (e.g. color out-of-date tables)

Click into details about your data such as its schema, when it was last updated, and the code that generated the data itself

Send feedback

- Collaborate with teammates
 - Create pipeline snapshots to share with other users

PREVIOUS

← [Code Repositories / Administer repositories / Compute Usage](#)

NEXT

[Navigation](#) →

© 2026 Palantir Technologies Inc. All rights reserved.

[Cookies Statement ↗](#)

[Privacy Statement ↗](#)

[Terms of Use ↗](#)

—[Cookies Settings](#)


[API Reference ↗](#)

[Send feedback](#)

Data connectivity & integration / Data Lineage / Overview

Data Lineage

Data Lineage is an interactive tool that facilitates a holistic view of how data flows through the Foundry platform.



With Data Lineage, you can:

- Easily find and discover datasets
 - Search to find datasets using Project, table, and column names
 - Click through Foundry Projects to browse data
- Explore [pipelines](#) through a powerful interface
 - Expand or hide ancestors and descendants of datasets

API Reference ↗ s of a group of tables at once

Visualize your pipeline through coloring (e.g. color out-of-date tables)

Click into details about your data such as its schema, when it was last updated, and the code that generated the data itself

Send feedback

- Collaborate with teammates
 - Create pipeline snapshots to share with other users

PREVIOUS

← [Code Repositories / Administer repositories / Compute Usage](#)

NEXT

[Navigation](#) →

© 2026 Palantir Technologies Inc. All rights reserved.

[Cookies Statement ↗](#)

[Privacy Statement ↗](#)

[Terms of Use ↗](#)

—[Cookies Settings](#)

[API Reference ↗](#)

[Send feedback](#)

Data connectivity & integration / Data Lineage / Navigation

Navigation


To make the best use of the Data Lineage application, you will need to know how to navigate the graphs, use tools, and configure branch and graph properties. The following numbered sections correspond to the numbers on the screenshot below:

1. [Lineage graph](#)
2. [Branch settings](#)
3. [Side panel](#)
 - [Search & Browse](#)
 - [Properties and Histogram](#)
 - [Manage Builds](#)
 - [Manage Schedules](#)
 - [Related Artifacts](#)
4. [Node details panel](#)
5. [Graph tools](#)
6. [Save graph](#)

API Reference ↗



Send feedback



Lineage graph

The graph is your workspace for arranging and manipulating nodes as you explore your data pipeline.

After adding nodes to the graph, you can add their related resources by clicking on the arrows on either side of the node or by using the **Expand** option in the [graph tools](#).

Nodes are arranged with auto-layout by default, but you can rearrange nodes manually by clicking and dragging them. To re-enable auto-layout, choose **Layout all nodes** under the **Layout** option in the [graph tools](#).

Click and drag to pan around the graph when in the default **Panning mode**. To use the cursor to select multiple nodes, switch to **Drag select** mode in the [graph tools](#) or hold **Shift** while clicking and dragging. You can select a node by clicking it, or select multiple

API Reference ↗ + click.

Send feedback

Branch settings

Select a branch from the list to explore data pipelines in that branch. The graph and any of the other helpers would show information based on the selected branch. If the branch does not exist for a resource, the listed fallback branches would be used instead (in the order they appear on the list).

To learn more about branching, see the [branching documentation](#).

Side panel

Search and browse

Use the search helper to find Foundry resources and add them to the graph. Use the free-text search or browse the tree to find resources. Add a resource by clicking on it or use the buttons at the bottom of the view to add all search results (including or excluding the content of sub-folders). Use the **Advanced** tab to add filters to your search and sort your results.

Warning

When viewing a folder with subfolders, you can recursively add *all* tables in all subfolders to the graph. Adding too many nodes at once may effect the graph's performance.

Properties and histogram

When you select a single node on the graph, the properties helper shows you the details of the resource. Depending on the type of resource you select, the properties helper shows available Foundry apps under the **Actions** menu and other links and actions (reporting issues, adding descriptions, etc.).

When you select multiple nodes on the graph, you will see the histogram helper. The helper displays API Reference ↗ properties and their values alongside the number of appearances of each value on the graph. By clicking on the values, the matching nodes are highlighted. If you want to drill down to just those resources, click on **Update selection**.

Send feedback

RESOURCE TYPE	# ▼
Code Workbook Dataset	36 / 36
Uploaded Dataset	16 / 16
Fusion Sync	4
Object Type	2
Writeback Dataset	1
TIME CREATED	# ▼
More than 30 days ago	52 / 57
TIME LAST BUILT	# ▼
More than 30 days ago	52 / 59

✕ Reset
✓ Update selection

i Use the **Copy names** button in the histogram to copy the names of all currently selected resources. The full names (including path) are copied to the clipboard as a comma-separated list.

Manage builds

The builds helper offers you three build strategies:

- Build only selected datasets
- Build all datasets between the selected datasets
- Build the selected datasets and all of their ancestors

[Learn more about managing builds.](#)

API Reference ↗

Manage schedules

Send feedback

The schedules helper allows you to set and edit build schedules for selected resources on the graph. [Learn more about build schedules.](#)



When viewing and creating schedules in Data Lineage, the schedules apply to the branches (including fallback branches) configured in the graph.

Related artifacts

The related artifacts helper displays artifacts directly linked to the nodes selected on the graph. Deleted and automatically saved files are excluded from the list unless chosen otherwise. You can also get to the same list of related artifacts by hovering over the right arrow of each node on the graph.

Node details

Click on a node to see more details:

- **Preview:** A sample of the data in the selected dataset.
- **History:** An overview of dataset change history. The overview includes tabs for logs, files, metadata, schema and job specifications.
- **Code:** If code was used to generate the dataset, it will display here
- **Data Health:** All the [health checks](#) set on the selected datasets.
- **Build timelines:** A Gantt chart of actual build time for the selected datasets.

Graph tools

The graph tools provide a set of graph exploration, navigation, and customization capabilities:

Node coloring

API Reference ↗

- [Expand](#)
- [Find](#)

Send feedback

- [Selection](#)

Node coloring

You can color the nodes on your lineage graph by several properties and metrics. Node coloring is commonly used to communicate lineage structure, troubleshoot issues, monitor pipelines health, and manage builds. You can also create your own custom coloring and arrange the graph based on the colors you assigned.

[Read more about node coloring options.](#)

i You can arrange your nodes on the graph by color group under **Layouts**.

Layout

The layout button provide various arrangement option for the nodes on the graph. **Layout all nodes** applies automatic layout for all the nodes on the graphs. When you select multiple nodes on the graph, you can apply other layouts (vertical, hierarchical, by level, etc.).

i You can use various useful keyboard shortcuts in the lineage graph. View the full list under the **Keyboard shortcuts** button at the top right corner of the app.

Expand

Use the **Expand** tool to expose ancestors and descendants of nodes in the graph. [Learn more about exploring data lineage.](#)

Find

API Reference ↗

Use **Find** to search for nodes on the graph. You can either search for tables or column names in datasets.

Send feedback

Selection

The **Selection** tool allows you to easily select nodes on the graph:

- **Select All:** Selects all the nodes currently on the graph.
- **Invert selection:** De-selects all currently selected nodes and selects the rest of the nodes on the graph.
- **Select children:** Adds all the direct children of the currently selected nodes to your selection
- **Select parents:** Adds all the direct parents of the currently selected nodes to your selection.

Save graph

You can save and share your lineage graph with other Foundry users in the following ways:

- **Save / Open:** Save your Data Lineage graph and re-open it by clicking on **Open graph**.
- **Get quick share link:** Generates a shareable link that provides read-only access to your graph.
- **Export graph to SVG:** Generates a static image of your lineage graph.

i Your branch choice is saved with your saved graph. If you load a graph with a different branch configuration than you currently have, you will be asked if you would like to switch branches to the saved branch configuration.

PREVIOUS
← Overview

NEXT
Branching data lineage →

© 2026 Palantir Technologies Inc. All rights reserved.

[API Reference ↗](#)

[Privacy Statement ↗](#)

[Terms of Use ↗](#)

Send feedback

Data connectivity & integration / Data Lineage / Branching data lineage

Branching data lineage

Data Lineage integrates with Global Branching, enabling you to view the state of your entities as they exist on a branch. You can select a global branch in Data Lineage to inspect branched resources, view branch-aware metadata, and search for entities on a branch.

For general information on Global Branching concepts and workflows, refer to the [Global Branching documentation](#).

i Data Lineage is not a branched resource. Graphs and graph configurations are not tracked as modified resources on a branch and cannot be merged back to `Main`. Data Lineage being branch-aware means that you can use it to view and interact with changes made on global branches.

Select a branch


To view resources on a global branch, use the branch selector in the Data Lineage toolbar and select the **Global branches** tab. Choose the branch you want to inspect.

When a global branch is selected, the lineage graph displays the state of resources as they exist on that branch, including any datasets, ontology entities, or other resources that have been added or modified.

[API Reference ↗](#)



[Send feedback](#)



- i** The branch selector also includes a **Dataset Branches** tab for selecting dataset-level branches. All global branches have a corresponding dataset branch, but not every dataset branch is tied to a global branch.

View branch-aware lineage


Once a global branch is selected, the lineage graph reflects the branch state across both datasets and table entities:

[API Reference ↗](#)

- **Datasets** display their branch-specific data, build status, and staleness

[Send feedback](#)

- **Ontology entities** such as object types, action types, and link types display their branch-specific metadata, including names and properties. Entities created or modified on the branch are reflected in the graph.
- **Links between resources** reflect the branch state, showing how datasets and ontology entities are connected on the branch.




i If your graph contains resources from multiple ontologies, only the ontology associated with the selected global branch will reflect branch-specific data. Resources from other ontologies will continue to show data from `Main`.

Search for entities on a branch

When a global branch is selected, the search in Data Lineage is branch-aware. You can search for ontology entities that exist on the branch, including entities that were created on the branch and do not yet exist on `Main`.

API Reference ↗

Send feedback




View entities modified on your branch

When a global branch is selected, the search panel displays a callout that allows you to add all entities modified on your branch to the graph with one click. This provides a quick way to visualize the full set of changes on your branch without searching for each entity individually.

The callout updates automatically as new resources are modified on your branch, so you can return to Data Lineage at any point during your branch workflow and see the latest set of changes.

[API Reference ↗](#)

[Send feedback](#)




Node details and previews

When you select a node on the graph while on a global branch, the details sidebar and preview panels display branch-specific information:

- **Dataset nodes** show branch-specific preview data, build history, and staleness.
- **Ontology entity nodes** show branch-specific metadata, including display names, descriptions, and properties as they exist on the branch.

[API Reference ↗](#)

[Send feedback](#)



Colorings and badges

Node [colorings](#) and badges in Data Lineage are branch-aware when a global branch is selected. Coloring options that depend on ontology metadata or dataset properties will reflect the state of the selected branch.

Saved graphs

When you save a graph while a global branch is selected, the global branch selection is saved with the graph configuration. Opening the saved graph restores the branch context, so you and other users will see the same branch-specific view of the lineage.

Known limitations

- **Dataset fallback branches:** When a global branch is selected, dataset fallback branches are disabled. The graph only displays data from the dataset branch associated with the [API Reference ↗](#) branch.
- **Single ontology per global branch:** Each global branch is associated with a single ontology. If your graph contains resources from multiple ontologies

Send feedback

to the selected global branch reflects branch-specific data. Resources from other ontologies continue to display data from **Main**.

- **Dataset-to-ontology links on older global branches:** Links between datasets and ontology entities (object types and link types) are indexed automatically for all entities on **Main** and for any global branch created on or after May 2026. Global branches created before May 2026 are not backfilled, and may show no dataset-to-ontology links for entities that have not been modified on the branch. To populate these links on an older branch, save a modification to the affected object type or link type on that branch to reindex its provenance.

PREVIOUS
← Navigation

NEXT
FAQ →

© 2026 Palantir Technologies Inc. All rights reserved.

[Cookies Statement ↗](#)

[Privacy Statement ↗](#)

[Terms of Use ↗](#)

— [Cookies Settings](#)

[API Reference ↗](#)

[Send feedback](#)

Data connectivity & integration / Data Lineage / FAQ

Data Lineage questions

The following are some frequently asked questions about Data Lineage.

For general information, view our [Data Lineage documentation](#).

- [How can I see the backing and writeback datasets for my object type in Data Lineage?](#)
- [What datasets in my pipeline also have a specific column?](#)
- [Who was the last person to modify a resource on this pipeline?](#)
- [How can I find which of my datasets have open transactions?](#)
- [Where are most of the datasets used in the pipeline stored?](#)
- [How can I share my unsaved Data Lineage graph?](#)
- [Why is my dataset not up-to-date?](#)

How can I see the backing and writeback datasets for my object type in Data Lineage?

- First, add your object to the Data Lineage graph by searching for it in the right panel (the tab with a magnifying glass icon). Select **Object types** to filter your search, then enter the name of the object for which you want to view the backing and writeback datasets.
- Next, select the arrow on the left side of your Object type to show its ancestors. This should produce one ancestor node if your object type is read-only and two ancestor nodes if your object type has writeback enabled. Make sure **Resource overview** is selected in the **Node color options** dropdown to see your Writeback Dataset colored as API Reference ↗ the top right. Backing schema dataset colors depend on the transform



Send feedback

- Your writeback and backing datasets for an object type will also have a small globe icon in the top right.

[Return to top](#)

What datasets in my pipeline also have a specific column?

1. First, ensure all desired datasets in your pipeline have been added to the Data Lineage graph.
2. Next, select desired datasets by using the **Select** mode in the **Tools** toggle in the upper left corner of the canvas.
3. Then open the **Histogram of selection properties** from the right side panel. Under the section titled **Frequent columns**, you will see the most frequent columns by column name in your selection.

Selecting one of these columns will highlight the datasets in your selection that contain this column.

[Return to top](#)

Who was the last person to modify a resource on this pipeline?

- First, ensure that all datasets of interest in your pipeline have been added to the Data Lineage graph.
- Next, select datasets by using Select mode from the **Tools** toggle in the upper left corner of your screen. Then, open the **Histogram of selection properties** from the right side panel.

[API Reference ↗](#) **Modified** section, you will see the last user(s) to modify datasets in your selection. Selecting a username will highlight the datasets that use the graph.

[Send feedback](#)

How can I find which of my datasets have open transactions?

In the dropdown menu in the top right side, choose **Build Status**. Now, you should be able to see if any dataset is currently running. Any such dataset has an open transaction.

Where are most of the datasets used in the pipeline stored?

- First, ensure that all datasets of interest in your pipeline have been added to the Data Lineage graph
- Next, select all datasets of interest with the **Select** mode from the **Tools** toggle in the upper left corner of your screen. Then, open the **Histogram of selection properties** from the right side panel.
- Under the section titled **Frequent folder paths**, you will see the most common folder paths for resources in your selection.

Selecting a golden path will highlight the resources in this path on the graph. Hovering over a folder path will show you the full path.

You can select multiple properties in the **Histogram of selection properties** panel such that the graph highlights all resources that satisfy your selection.

How can I share my unsaved Data Lineage?

Send feedback

To share your unsaved Data Lineage, select the arrow in the top right corner near Save. Once there, you can see a quick share link.

[Return to top](#)

Why is my dataset not up-to-date?

There are a few reasons why your dataset may not be up-to-date.

Consider the following reasons why your dataset may not be up-to-date:

- Is your dataset build failing?
- Is there an upstream dataset that has not built and is not up-to-date?
- Have you received up-to-date data from the source?

You can easily answer these questions in Data Lineage:

1. First, verify the status of each of the resources in your pipeline by opening up the dataset of interest in Data Lineage and then right-clicking on the node.
2. Then, select **Expand node....** You can view all ancestor nodes for that dataset by selecting the double left arrow above **Expand parents....**
3. Next, select the **Build status** option in the **Node color options** dropdown menu in the top right to view the build status of every resource in your pipeline. This view of your pipeline will make it easier to diagnose stale datasets.

[Return to top](#)

PREVIOUS
← [Branching data lineage](#)

NEXT
[Graphs / Explore data lineage](#) →

[API Reference ↗](#)

Technologies Inc. All rights reserved.

[Cookies Statement ↗](#)

[Privacy Statement ↗](#)

[Send feedback](#)

[Terms of Use ↗](#)

— [Cookies Settings](#)


[API Reference ↗](#)

[Send feedback](#)

Explore data lineage

Data lineage helps you understand how your data came to be. There are various ways to explore data pipelines in the Data Lineage app. Consider one common path:

1. Using the **Search** helper, find your resource (for example, a dataset or an object type) and add it to the graph.
2. Click on the left arrow of the node to expose the direct parents of the resource.




3. To expand your graph, select the next resource on the graph and click the **Expand** button in the graph tools.

4. API Reference ↗ Iron button to define the number of levels to expose. Click the double- to expand all the way to the raw data (or all the way to the first dataset in the lineage).



Send feedback

- i** Adding too many nodes simultaneously may affect the graph's performance and usability. Keep a manageable number of nodes by checking the node count in the **Expand** tool.




- i** You can find the relation between two nodes on the graph by selecting the **Expand** button and adding all nodes in between the resources or all common ancestors/descendants.

5. Get more information on one of the datasets by selecting the dataset and using the bottom panel to display a preview of the data.
6. Click on **Code** to view how the dataset was created.

[API Reference ↗](#)

[Send feedback](#)



7. Click on **View in code workbook** or **View in repository** to see the original code and make changes as needed (subject to permissions).

i Some options may be unavailable for some datasets depending on the type of resource. For example, **Code** is only available for Code Workbook or Code Repositories. For Fusion sheet syncs with no code to show, you may have the option to view the source sheet and make your changes there (if you have appropriate permissions).

PREVIOUS
← [FAQ](#)

NEXT
[Explore artifacts and ontology entities](#) →

© 2026 Palantir Technologies Inc. All rights reserved.

[API Reference](#) ↗

[Privacy Statement](#) ↗

[Terms of Use](#) ↗

[Send feedback](#)

Data connectivity & integration / Data Lineage / Graphs / Explore artifacts and ontology entities

Explore artifacts and ontology entities

You can find Foundry artifacts and ontology entities related to your datasets in Data Lineage. The Data Lineage interface allows you to navigate directly to these resources and see how they fit in your ontology.

Find related artifacts

In your data lineage graph, select a dataset. Then, select **Related items** in the right sidebar to expand the **Related artifacts** panel. The **Related items** icon will show a badge with the number of artifacts related to the selected dataset. In the artifacts panel, you can see a list of related resources throughout Foundry, including Contour visualizations and Slate applications.

API Reference ↗



Send feedback

The screenshot displays a data lineage tool interface. On the left, a graph shows data flows between nodes like '[Internal Training] Airport', 'delays_edited', '[Internal Training] Airline', and 'Airline Comparison Bubble Plot'. A 'Related items' tooltip is visible over the graph. On the right, the 'Related artifacts' sidebar is active, showing a list of artifacts with filters and sorting options. The artifacts listed are:

- Airline Comparison Bubble Plot**: Created Mar 21, 2022, 7:15 AM; Last modified Apr 11, 2022, 6:39 AM.
- Embedded Contour (Airlines)**: Created Mar 21, 2022, 7:02 AM; Last modified Apr 11, 2022, 6:45 AM.
- Embedded Contour (Airlines)**: Created Oct 25, 2021, 8:06 AM; Last modified Nov 23, 2021, 9:48 AM.
- Slate: Sample Combined Data**: Created Oct 12, 2021, 3:08 PM; Last modified Oct 25, 2021, 9:17 AM.
- Contour: Data-oriented Embed**: Created Oct 12, 2021, 12:36 PM; Last modified Oct 12, 2021, 2:11 PM.
- Sample Embedded Report**: Created Oct 5, 2021, 3:01 PM; Last modified Oct 25, 2021, 12:55 PM.


Click on the node icon next to a resource to zoom in on the related dataset, or click the resource to open it in the corresponding application in a new tab. You can filter the list of related artifacts to include different item types and sort the list by oldest, newest, name, path, or last modified.

Find ontology entities

Find object types defined by datasets in your lineage graph by selecting the dataset and opening the **View node properties** panel in the right sidebar.

[API Reference ↗](#)

[Send feedback](#)




In the **About** tab, you will see any object types that were created with the selected dataset. Click the **Settings** icon next to the object type to view its configuration in a new Ontology manager tab.



You can also add object types to your data lineage with the **Search Foundry** tool in the right sidebar. Use a basic or advanced search to find an object type and select it from the list to add it to your graph. You can then view link types related to the object type and use the graph to visualize connections between your datasets and the newly added object type.

API Reference ↗

Send feedback



[Learn more about creating an Ontology.](#)

PREVIOUS
← Explore data lineage

NEXT
Save and share a graph →

© 2026 Palantir Technologies Inc. All rights reserved.

[Cookies Statement ↗](#)

[Privacy Statement ↗](#)

[Terms of Use ↗](#)


—Cookies Settings

[API Reference ↗](#)

[Send feedback](#)

Save and share a graph

Data Lineage allows you to easily save your graph and share it with other users. You can find multiple ways to save and share by selecting the **Actions** tab in the upper right of the application and selecting a method from the dropdown menu.



Save a graph


Select the following options in the **Actions** dropdown menu to save your Data Lineage graph:

- **Save:** Save your Data Lineage graph to where it currently lives in your files or Project.
- **Save as...:** Choose a name for your lineage graph and save it to a new location in your file

API Reference ↗




Send feedback




You can also open a previously saved graph with the following option:

- **Open graph:** Choose to open a different saved graph to which you have access, or open the **Clipboard** tab to enter the resource identifier (RID) of a dataset, schedule, graph, or path.



Send feedback

- i** Your branch choice is saved with your saved graph. If you load a graph with a different branch configuration than you currently have, you will be asked if you would like to switch branches to the saved branch configuration.



Share a graph

You can share a graph with other users using the options below:

- **Get quick share link:** Generate a shareable link that provides read-only access to your graph. Note that this option is only available for users belonging to the same Organization. To share a graph across Organizations, ensure the graph is saved in [a shared Project accessible to both Organizations](#).
- **Export graph to SVG:** Generate and download a static image of your lineage graph in .svg format.

You can also select **Share** in the upper right of the application to open the sidebar and view **Roles** details. From here, you can turn on link sharing or give a user or group access to your graph.

API Reference ↗

Send feedback

Aviation Pipeline

Details > Access > Roles

Link sharing

https://

Viewer

People who visit the link will be granted the **Viewer** role. They will still need to meet access requirements to be able to access the file.

Roles

Add a user or group...

Node coloring

There are several built-in options for coloring graph nodes to give you more information about your pipeline:

Coloring option	Description
No color	Would remove coloring altogether
Custom color	Allows you to select nodes and assign them a color by clicking on the Color button
Data Catalog	Nodes would be colored based on the Data Catalog collection they are in. If the node belongs to more than one collection it would be colored as “In multiple collections”
Folder	Colors the nodes by name of the folder the resource is located in
Issues	Colors the nodes by the number of Foundry issues assigned to them. This option would also allow you to filter by issues labels.
Permissions	Colors the nodes by the level of access the user has to the data or the resource. If you have access to the resources on the graph, this view also allows you to choose any Foundry user and view their permissions.
Project	Colors the nodes based on the Foundry Project they are in.
Repository	Colors the nodes based on the code repository used to create them. You can either color the nodes by the name of the repository, or by its type (e.g. Code Repository, Code Workbook).

API Reference ↗

This view colors the node by details of the resource. The details of the resource generally refer to the way the resource was created (such as in Contour, Code Workbook, Fusion spreadsheet etc.).

Send feedback

Coloring option	Description
Resource type	Colors the nodes by Foundry resource types.
Build status	Indicates the current build status of each dataset on the graph. If the nodes are grouped the more severe status would be presented.
Data Health	Indicates the status of resources health checks with the ability to filter to only watched health checks. If the nodes are grouped, the color of the group would indicate the most severe health check status of the group.
Out-of-date	This option would indicate if the data or logic is out of date relative to the dataset ancestors. Out-of-date with parent means a direct parent of the resource had been updated and the resource itself hasn't yet updated accordingly. Out-of-date with ancestor means the resource is up-to-date with its direct parents, but there is a resource upstream that is more updated. This options allows you to filter for two types of updates: Data and Logic. Data out-of-date means the data was updated in an ancestor and the resource hasn't yet picked up the update in a build. Logic out-of-date means job-specs has changed.
Schedule count	Indicates the amount of build schedules set on a dataset with the option to filter out paused schedules.
Sync status	If there are syncs set on the dataset, this option would indicate the status of the sync
Time last built	Indicates the time since the last time the dataset was successfully built.
Build duration	According to the most recent successful build of each resource, this option would indicate the approximate build duration
Files	Colors the nodes on the graph by files-related metrics: Average file size, count of files and dataset size
API Reference ↗	Colors nodes by the number of rows in each dataset. If row count does not exist, it could be calculated in the dataset details helper or in the dataset view in Foundry (dataset app).
Spark usage	Colors each node by Executor run/CPU time in a job.

Send feedback

Coloring option	Description
User views	Colors nodes by the number of user views
Branch	Indicates the currently viewed branch of each node on the graph.
Code Status	Indicates the code status for this node/dataset. CI running means CI checks are currently running for this node. CI Failed means that CI checks failed on this node. Out of date means that the code is out of date for this node. Unavailable means that the node/dataset is not a stemma backend or that the user is lacking permissions.
Storage	Indicates where data is stored. Will be Foundry unless you are using Virtual Tables .
Compute	Indicates the compute engine used by each node on the graph. For transforms run in Foundry, this will show the type of compute used (Spark or Flink, for example). For transforms that use compute pushdown, this will indicate the external compute engine used (BigQuery, Databricks or Snowflake, for example).
Transaction type	Indicates each nodes transaction type: Append or Snapshot

[PREVIOUS](#)
 ← Save and share a graph

NEXT
 Graph elements reference →

© 2026 Palantir Technologies Inc. All rights reserved.

[Cookies Statement ↗](#)

[Privacy Statement ↗](#)

[Terms of Use ↗](#)



[API Reference ↗](#)

[Send feedback](#)

Data connectivity & integration / Data Lineage / Graphs / Graph elements reference

Graph elements reference

Node types

Node	Type	Description
	Data source	This is the name of the data source as it appears in Data Connection . Learn more about the different source types.
	Dataset	Foundry datasets and the lineage between them. The color of the dataset node depends on user selection . Dashed border indicates unstructured datasets.
	Object type	Ontology object types . The icon and color of the node depends on the definition of each object type. When clicking on the “link” icon next to the object type name, Data Lineage shows the relations between this object type and other object types.
	Artifact	Data Lineage exposes different Foundry artifacts like: Contour analyses, Reports , etc. The color of the node depends on the artifact type, which is indicated at the top of the node.

Node indicators appear on top of dataset nodes and provide additional information about the resource.

Indicator	Type	Description
	Open issues	This indicator signals there are currently open issues associated with the node on the graph. Hovering over this signal indicates the number of open issues.
	Defines an object type	This indicator appears on datasets that are used to define Ontology object types. Hovering over the right arrow of allows you to expose those linked object types. Learn more about object types.
	Syncs	Datasets with this indicator on them have syncs to other databases or systems. You can view these syncs by selecting the node and opening the Properties panel, or by opening the “Details” tab in Dataset Preview (right click on the node and click on Open).
	Trashed	This indicator appears on nodes representing deleted datasets or artifacts. Deleted nodes are also partially faded with their name crossed out.

Data connectivity & integration / Data Lineage / Understand and manage datasets / View dataset preview and logic

Preview and logic

The Data Lineage interface allows you to view previews of selected datasets or media sets, as well as examine the associated code to understand the logic behind the dataset or media set.

Preview

To see a preview of a dataset or media set, select it in your data lineage graph, then choose the **Preview** tab in the bottom left of the interface.

Media set


When the media set preview expands, you can view the contents of your media set. [Learn more about media sets.](#)

Example of PDF preview:


API Reference ↗



Send feedback



Example of audio preview:



[API Reference ↗](#)

Dataset

[Send feedback](#)

When the dataset preview expands, you can scroll through the first 300 rows of the selected dataset. You can also search for specific columns using the **Search columns...** field to the right of the preview window. The preview of your dataset will look different depending on the type of data within your dataset.

Preview

History

</> Code

Build timeline

Data health

1 node selected

flights

Showing 300 of 4.2m rows61 columns

Q

Search columns...

	op_carrier String	tail_num String	op_carrier_fl_num Integer	origin String	origin_city_name String	origin_state_abr String	origin_state_nm String	dest String	dest_city_name String
1	OO	N432SW	4712	DTW	Detroit, MI	MI	Michigan	CAK	Akron, OH
2	YX	N731YX	3651	CLT	Charlotte, NC	NC	North Carolina	EWR	Newark, NJ
3	F9	N702FR	1262	LAS	Las Vegas, NV	NV	Nevada	ORD	Chicago, IL
4	AS	N403AS	450	SEA	Seattle, WA	WA	Washington	ONT	Ontario, CA
5	OO	N138SY	5793	SAN	San Diego, CA	CA	California	LAX	Los Angeles, CA
6	AA	N568UW	1916	CLT	Charlotte, NC	NC	North Carolina	ORD	Chicago, IL
7	B6	N599JB	2074	TPA	Tampa, FL	FL	Florida	EWR	Newark, NJ
8	AS	N525VA	1598	SFO	San Francisco, CA	CA	California	PSP	Palm Springs, CA
9	OO	N947SW	5157	SPI	Springfield, IL	IL	Illinois	ORD	Chicago, IL
10	YX	N882RW	5885	LGA	New York, NY	NY	New York	CHS	Charleston, SC
11	UA	N19141	341	EWR	Newark, NJ	NJ	New Jersey	DEN	Denver, CO
12	DL	N312DN	2137	TPA	Tampa, FL	FL	Florida	ATL	Atlanta, GA
13	AA	N941NN	1097	SEA	Seattle, WA	WA	Washington	ORD	Chicago, IL
14	UA	N69847	599	BOS	Boston, MA	MA	Massachusetts	ORD	Chicago, IL
15	B6	N958JB	51	BOS	Boston, MA	MA	Massachusetts	MCO	Orlando, FL

Logic

Select the **Code** tab to view the code logic of the selected dataset or media set. From the **Code** view, you can make quick edits, search for items, or open the code in the repository or other application used to derive the data.

PreviewHistoryCodeBuild timelineData health1 node selected

Quick editFind...View in repository

```
1 from transforms.api import transform_df, Input, Output, configure
2 from pyspark.sql import functions as F, types as T
3
4 from datetime import datetime
5 from dateutil.relativedelta import relativedelta
6 import pytz
7 from timezonefinder import timezonefinder
8 import pandas as pd
9 from itertools import chain
10
11
12 @configure(profile=['EXECUTOR_MEMORY_MEDIUM'])
13 @transform_df(
14     Output("/Public/Foundry Training and Resources/Foundry Reference Project/Datasource Project: Flight Control System/flights/enriched/flights"), # noqa
15     flights=Input("/Public/Foundry Training and Resources/Foundry Reference Project/Datasource Project: Flight Control System/flights/processed/flights"), # no
16     carrier_lookup=Input("/Public/Foundry Training and Resources/Foundry Reference Project/Datasource Project: Flight Control System/flights/clean/carrier_looku
17     cancellation_code_lookup=Input("/Public/Foundry Training and Resources/Foundry Reference Project/Datasource Project: Flight Control System/flights/clean/can
18     distance_group_lookup=Input("/Public/Foundry Training and Resources/Foundry Reference Project/Datasource Project: Flight Control System/flights/clean/distan
19     delay_groups_lookup=Input("/Public/Foundry Training and Resources/Foundry Reference Project/Datasource Project: Flight Control System/flights/clean/delay_g
```

Uploaded and writeback datasets do not have associated code to view in Data Lineage.

API Reference ↗

© 2026 Palantir Technologies Inc. All rights reserved.

[Cookies Statement ↗](#)

[Privacy Statement ↗](#)

[Terms of Use ↗](#)

— [Cookies Settings](#)

[API Reference ↗](#)


[Send feedback](#)

Data connectivity & integration / Data Lineage / Understand and manage datasets / View build timeline

View build timeline

Use the **Build timeline** tool in Data Lineage to view the build history of your datasets.

In Data Lineage, click on **Build timeline** in the bottom left of your window. This action expands the view panel to display a Gantt chart of builds that took place during the period of time of your choice. You can select the number of days or hours you would like to view in your timeline, ranging from one hour to ten days. You can also choose to display the builds by color based on schedule or job status.




To view the build timeline of a specific dataset, select the dataset on your graph. Select multiple datasets with the **Drag select mode** tool or by holding **Ctrl / Command** while clicking.

To view the build timeline of a specific dataset, select the dataset on your graph. Select multiple datasets with the **Drag select mode** tool or by holding **Ctrl / Command** while clicking.

[API Reference ↗](#)

Send feedback



See more details about the build, job, and schedule by clicking on the links within the job information window.

PREVIOUS
← View dataset preview and logic

Understand out-of-date datasets → NEXT

© 2026 Palantir Technologies Inc. All rights reserved.

[Cookies Statement ↗](#)

[Privacy Statement ↗](#)

[Terms of Use ↗](#)

— [Cookies Settings](#)

[API Reference ↗](#)

[Send feedback](#)

Data connectivity & integration / Data Lineage / Understand and manage datasets / Understand out-of-date datasets

Understand out-of-date datasets

There are a few reasons why your dataset may not be up to date. Common scenarios to explore are:

- Is my dataset build failing?
- Is there an upstream dataset that hasn't built and isn't up to date?
- Have we received up-to-date data from the source?


You can easily answer these questions by using Data Lineage.

- First, verify the status of each of the resources in your pipeline by opening up the dataset of interest in Data Lineage and right-clicking on the node.


API Reference ↗



Send feedback




- Then, select **Expand node**. You can see all of the ancestor nodes for that dataset by clicking the double left arrow above **Expand parents**.



Send feedback

- Next, select the **Build status** option in the **Node color options** dropdown in the top right of Data Lineage to see the build status of every resource in your pipeline. This view of your pipeline will make it much easier to diagnose stale datasets.



PREVIOUS

← [View build timeline](#)

NEXT

[Find datasets with a given column](#) →

© 2026 Palantir Technologies Inc. All rights reserved.

[Cookies Statement](#) ↗

[Privacy Statement](#) ↗

[Terms of Use](#) ↗

— [Cookies Settings](#)

[API Reference](#) ↗


[Send feedback](#)

Data connectivity & integration / Data Lineage / Understand and manage datasets / Find datasets with a given column

Find datasets with a given column

You can easily search for specific dataset columns within your Data Lineage graph:

- First, ensure you added all datasets of interest in your pipeline to your lineage graph.
- Next, select all datasets of interest by using **Drag select mode** in the Tools toggle in the upper left hand corner of the app. You can also hold down `Ctrl / Command` to select multiple nodes at once, or use `Ctrl / Command + A` to select all nodes.




- Then, select **View histogram of selection properties** from the Data Lineage sidebar.


[API Reference ↗](#)



[Send feedback](#)



- Under the **Frequent Columns** section, you can see the most frequent columns by name in your selection.
- Click one of the columns to highlight the datasets in your selection that contain this column.



PREVIOUS
← Understand out-of-date datasets

NEXT
Build datasets →

[API Reference ↗](#)

© 2020 Palantir Technologies Inc. All rights reserved.

[Cookies Statement ↗](#)

[Send feedback](#)

[Privacy Statement ↗](#)

[Terms of Use ↗](#)

— [Cookies Settings](#)

[API Reference ↗](#)

[Send feedback](#)

Data connectivity & integration / Data Lineage / Understand and manage datasets / Build datasets

Build datasets

You can use the Data Lineage graph to see which datasets in your pipeline are out of date, and then use the Builds helper to start builds directly from Data Lineage.

i Builds triggered from Data Lineage always apply to the branches (including fallback branches) configured in the graph.

The following are a few common build workflows:

- [Build all ancestors](#)
- [All transforms in between selected datasets](#)
- [Selected dataset\(s\) only](#)

Build All Ancestors

This strategy builds the selected datasets and all ancestor datasets, to ensure that the selected datasets become completely up to date.

i By default, this builds only ancestors that are out of date, but you can choose to force a re-build of up-to-date datasets. Forcing a re-build can be expensive in terms of build time and resources.

[API Reference ↗](#)

sets to the graph or open a saved snapshot.
e dataset that you want to build.

Send feedback

3. In the Builds helper, choose **All ancestor datasets**, then click **Next**.

i Clicking **Next** will *not* trigger any builds yet. You will simply see a preview of the datasets to be built.


4. If you want to force a re-build of up-to date datasets, click **Force build** on up-to-date datasets.

5. After examining the list of datasets to be built, click **Run build** to trigger the builds.

i If you decide you do not want to build *all* out-of-date ancestors, you must click **Cancel** on the current build preview, then change the nodes you have selected. You cannot change your selection from the build preview screen.

[API Reference ↗](#)

[Send feedback](#)



All transforms in between selected datasets

This strategy lets you bind your builds to a subset of your pipeline. A common use case for this strategy can occur when new raw data regularly lands in your pipeline and there is a particular dataset that you want to update to reflect the new data, but you don't want to build *all* out-of-date ancestors. You can then use Data Lineage to determine which other datasets need to be built to bring your dataset of interest more up to date.

1. Add the dataset you ultimately want to build to the graph.
2. Add any raw datasets to the graph (or any upstream dataset)
3. Select all nodes.
4. In the Builds helper, choose the **All transforms in between selected dataset(s)** strategy, then click **Next**.

API Reference ↗

Clicking **Next** will not trigger any builds yet. You will simply see a preview of the build based on the nodes you have selected. You can now see exactly what needs to be built to update your dataset of interest. You may

Send feedback

all datasets – maybe there is a very large derived dataset that should only build once a day – so click **Add all to graph** at the bottom of the list.

Selected Datasets

This strategy allows you to pick individual datasets that you want to build. If there are dependencies between the datasets, builds would be executed in the right order to assure descendants are built after their ancestors were built.

i If you want to change the datasets you are building, you must click **Cancel** on the current build preview, change the nodes you have selected, then enter a new preview. You cannot change your build selection from the build preview screen.

After examining the final list of datasets to be built, click **Run** build to trigger the builds.

PREVIOUS
← Find datasets with a given column

NEXT
Manage schedules →

© 2026 Palantir Technologies Inc. All rights reserved.

[Cookies Statement ↗](#)

[Privacy Statement ↗](#)

[Terms of Use ↗](#)

—Cookies Settings


[API Reference ↗](#)

[Send feedback](#)

Data connectivity & integration / Data Lineage / Understand and manage datasets / Manage schedules

Manage schedules

Data Lineage allows you to easily manage build schedules within your lineage graph. In the right sidebar, select **Manage schedules** to open the schedule details pane.




You will see the schedules related to selected datasets in your graph. Click on a schedule to see more details:

[API Reference ↗](#)



[Send feedback](#)



- **Latest run:** The status of the latest run of the schedule.
- **Last update:** A timestamp of when the last update took place and the user who made changes
- **Target datasets:** A list of downstream datasets including in the build schedule.
- **When to build:** Displays the build schedule trigger determined when creating the build schedule. For example, a build schedule can be set to run **when specific datasets update**.
- **Build scope:** Defines the Project or user datasets included in the build and the permissions used to run the build.

[API Reference](#)
[Scheduling builds in the Building pipelines documentation.](#)

Send feedback

PREVIOUS

NEXT

© 2026 Palantir Technologies Inc. All rights reserved.

[Cookies Statement ↗](#)

[Privacy Statement ↗](#)

[Terms of Use ↗](#)

—[Cookies Settings](#)

[API Reference ↗](#)

Send feedback

Data connectivity & integration / Data Lineage / Understand and manage datasets / Roll back a pipeline

Roll back a pipeline [Planned deprecation]

⚠ Planned deprecation

Rolling back a pipeline is in the [planned deprecation](#) phase of development and will be unavailable for use after November 30th. We recommend using the [roll back a dataset](#) feature instead.

When building your pipeline, you may need to roll back a [dataset](#) and all of its downstream dependents to an earlier version. There can be many reasons for this, including the following:

- You identified a mistake in the logic required to build a dataset and need to revert it.
- Incorrect data was pushed into your pipeline from an upstream source.
- An outage occurred, and you want to quickly navigate back to an earlier state of your pipeline.

The pipeline rollback feature allows you to revert back to a [transaction](#) of an upstream dataset. When performing a rollback, the data provenance of the upstream dataset transaction is used to identify its downstream datasets and their corresponding transactions to create a final pipeline rollback state. Typically, this process would require several steps to properly roll back each affected dataset. With pipeline rollback, this is reduced to a few simple steps discussed below, along with the ability to preview the final pipeline state before confirming and proceeding with the rollback. Pipeline rollback also ensures that the incrementality of your pipeline is preserved.

API Reference ↗




Send feedback

As you set up your rollback, you can choose to exclude any downstream datasets; these datasets will remain unchanged as the pipeline is rolled back to the selected transaction.

⚠ If a dataset backs an [object type](#) stored using Object Storage v2, manual intervention is required to ensure that the object type is [reindexed](#) with a successful run of the [replacement pipeline](#) to reflect the state after the rollback.

Execute a pipeline rollback


1. Navigate to a [Data Lineage](#) graph containing the upstream dataset you would like to roll back.
2. Select the dataset in the graph. Then, from the branch selector at the top of the graph, select the branch on which you would like to perform the rollback.
3. Select **View node properties** in the panel on the right.



4. Select **Actions**, then **Rollback**.
5. Under **Selected transaction**, choose the transaction to which you would like to roll back.

[API Reference ↗](#)

[Send feedback](#)



After choosing the transaction, downstream datasets will automatically be found and the states they will revert to if the rollback is actioned will be displayed.

Resource types that are unable to be rolled back, including streaming datasets, media sets, and restricted views, will be displayed under the **unsupported resources** section. Transactional datasets on which you do not have `Edit` access will also be included in this list.

6. Select the timestamp under each dataset to navigate to the [History page](#) of the input, where the corresponding transaction will be highlighted.

[API Reference ↗](#)

[Send feedback](#)

Exit rollback

Only transactional datasets are supported for rollbacks. Any other downstream resources, including media sets, streaming datasets, or virtual tables will remain unchanged. [See documentation](#)

Selected dataset to rollback

Selected transaction

Wed, Feb 5, 2025, 11:32 PM

Select

Downstream Rollback ⓘ

6 unsupported resources

Wed, Feb 5, 11:33 PM

Wed, Feb 5, 11:34 PM

Wed, Feb 5, 11:48 PM

Wed, Feb 5, 11:32 PM

Wed, Feb 5, 11:33 PM

Datasets excluded from rollback ⓘ


Remove nodes from auto-rollback to keep current.

Rollback


7. Select any datasets to exclude from the rollback by selecting — to the right of the dataset name. Once excluded from rollback, the dataset will appear in the **Datasets excluded from rollback** section.

API Reference ↗

Send feedback



8. To add an excluded dataset back to the rollback, select **+** to the right of the dataset name.



9. After finalizing the state of your desired rollback, select **Rollback**. A confirmation dialog will appear.

Confirm rollback

You may not have permissions to discover all downstream outputs of a dataset. Rolling back may not update the outputs to which you do not have access, resulting in failing builds, unexpected transactions (such as snapshot instead of incremental), or other consequences.

This action will rollback transactions on **5 datasets** and reset incremental state on **2 datasets**.

To confirm rollback, enter the branch name: **master**

master

Cancel

Confirm rollback

10. Enter the branch name as confirmation, then select **Confirm rollback** to proceed.

← Exit rollback

Rolled back 7 datasets

 even_duplicate Wed, Feb 5, 11:33 PM	✓
 even Wed, Feb 5, 11:34 PM	✓
 adds_timestamp_duplicate Wed, Feb 5, 11:48 PM	✓
 input Wed, Feb 5, 11:32 PM	✓
 adds_timestamp Wed, Feb 5, 11:33 PM	✓
 media_set_downstream	✓
 test_output	✓

Rollback successful

Exit

11. Once the rollback is complete, navigate to the **History** tab of the datasets and notice that the rolled back transaction is now crossed out, as shown below:

Incremental_input >

incremental_output

PreviewSQL scratchpadHistoryCodeBuild timelineData health

Summary

Job duration

Jan 12, 4:35 AM
Delete - Foundry

Jan 11, 5:50 AM
Delete - Foundry

Date range:Jun 12, 2025Jun 19, 2025


1 node selected

9 Total jobs9 Successful9 Failed9 Cancelled-- Median duration


Show running average

Depending on how your pipeline changed between the current state and the rollback transaction, you may receive one of the warnings shown below on datasets selected for downstream rollback:

- If the rollback transaction was built using a different JobSpec (logic) than the latest transaction, the transaction will show the following warning: `The dataset logic was changed after the selected transaction time. Rolling back this dataset will not revert logic changes.` If you receive this warning, you may need to inspect how the logic changed over time and determine if you can apply the logic again after the rollback and before any new dataset builds.



- If a downstream dataset did not exist at the time of the selected rollback transaction, the dataset will show the following warning: `This dataset did not exist at the time of the transaction selected for rollback. Its incremental state will be reset, though transactions and data will be unchanged after rollback. It will run as a snapshot in the next build.` If you receive this warning, note that the rollback will reset the dataset's job history as if it was never built. The next build on the dataset will not be run incrementally (if the dataset is an incremental dataset) since it will be treated as a completely new dataset being built for the first time.



[API Reference ↗](#)

Considerations and limitations

[Send feedback](#)

When rolling back a pipeline, keep the following considerations in mind:

- Only transactional datasets are supported for rollbacks. Any other downstream resources, including media sets, streaming datasets, or virtual tables will remain unchanged.
 - These resources will be displayed under the **unsupported resources** section of the rollback preview panel.
- You are only able to roll back to a *successful* transaction.
- It is not possible to roll back to a [retentioned transaction](#)
 - However, you can roll back to a **DELETE** retention transaction.
- You can only roll back datasets to which you have access. If you do not have access to certain downstream datasets, they **will not** be rolled back.
 - It is possible to have **Edit** access on an upstream dataset that is then used to build a downstream dataset to which you only have **View** access. In this case, the downstream dataset will be considered an "unsupported resource".
- Only transactional datasets with a [JobSpec](#) can be rolled back; uploaded datasets are not supported for rollback.
- A maximum of **50 transactions** can be rolled back on the upstream dataset at a given time.

PREVIOUS

← [Manage schedules](#)

NEXT

[Roll back a dataset](#) →

© 2026 Palantir Technologies Inc. All rights reserved.

[Cookies Statement](#) ↗

[Privacy Statement](#) ↗

[Terms of Use](#) ↗

— [Cookies Settings](#)

[API Reference](#) ↗

[Send feedback](#)

Data connectivity & integration / Data Lineage / Understand and manage datasets / Roll back a dataset

Roll back a dataset

When building a pipeline, you may need to roll back a [dataset](#) to an earlier version. There may be various reasons for this, including the following:

- You identified a mistake in the logic required to build a dataset and need to revert it.
- Incorrect data was pushed into your pipeline from an upstream source.
- An outage occurred, and you want to quickly navigate back to an earlier state of your dataset.

The dataset rollback feature allows you to update the data and job history of a dataset. If the dataset is being built [incrementally](#), the dataset rollback feature also ensures that the incrementality of your dataset is preserved.

Types of dataset rollbacks

The two types of possible rollbacks on a dataset are as follows:

1. **[Rolling back to an earlier transaction](#)**: Performed when there is a previous transaction to roll the dataset back to.
2. **[Forcing a snapshot on the next build](#)**: Typically applicable for incremental workflows when there is no previous transaction to roll back to, meaning that the dataset needs to be rolled back to a state before the branch or the dataset was created.

API Reference ↗

...ally force a snapshot on the next build of a dataset, but you intended to roll back to an earlier transaction, do not proceed with a rollback, as this could leave the dataset in a partially rolled back state.

Send feedback

Instead, build the dataset; it will run as a snapshot since the dataset was configured to snapshot on the next build, and then carry out the intended rollback.

Considerations and limitations

When rolling back a dataset, keep the following considerations in mind:

- Only transactional datasets are supported for rollbacks.
- You are only able to roll back to a *successful* transaction.
- It is not possible to roll back to a transaction that was deleted based on a [retention policy](#). However, you can roll back to a transaction that was deleted by a dataset rollback in Data Lineage.
- You can only roll back a dataset on which you have the [Editor role](#).
- After a rollback is carried out, the logic backing the dataset will be left unchanged and will need to be updated in order to apply to the next build.
- If the branch on which a rollback is being performed does not exist on the dataset, the rollback will be applied to a fallback branch.

 If multiple outputs are built together by the same job, each output would need to be rolled back to a transaction that was built by the same previous job.


Similarly, if multiple outputs are built together, all of them would need to be forced to run as a snapshot on the next build.

Rolling back or requesting a snapshot for only a subset of outputs would result in outputs being left in an inconsistent state for subsequent jobs. Specifically, for incremental workflows:

- Partial rollbacks could cause the subsequent job to run unexpectedly as a snapshot.
- Partial force snapshot requests would cause the subsequent job to continue running incrementally.

[API Reference ↗](#)

[Send feedback](#)



Roll back to an earlier transaction

1. Navigate to a [Data Lineage](#) graph containing the dataset you would like to roll back.
2. Select the dataset in the graph. Then, from the branch selector at the top of the graph, select the branch on which you would like to perform the rollback.



3. Select the **History** tab in the panel at the bottom of the page.
4. Select the transaction to roll back to.

[API Reference ↗](#)

[Send feedback](#)

dataset

Preview

SQL scratchpad

History

Code

Build timeline

Data health

Summary

Overview

Logs

Files

Metadata

Schema

Job spec

Advanced

Feb 16, 3:13 PM

Update • scheduler

Part of build • 1 job

(54s)

Feb 16, 3:12 PM

Update • scheduler

Part of build • 1 job

(52s)

Feb 16, 3:11 PM

Update • scheduler

Part of build • 1 job

(56s)

Feb 16, 3:10 PM

Update • scheduler

Part of build • 1 job

(53s)

Feb 16, 3:09 PM

Update • scheduler

Part of build • 1 job

(52s)

Feb 16, 3:08 PM

Update • scheduler

Part of build • 1 job

(56s)

Transaction details

Status

Committed

Type

Update

Duration

55s

Start time

Feb 16, 2025, 3:09 PM

Close time

Feb 16, 2025, 3:10 PM

Files

2

Size of files

625B

Started by

Foundry

Transaction RID

fe-b523-2ebbc0aa2bbc

Rollback

Rollback to transaction

Job details

Status

Succeeded (Part of a build)

Type

transforms python-1

Duration

53s (part of a 57s build)

Start time

Feb 16, 2025, 3:09 PM

End time

Feb 16, 2025, 3:10 PM

Started by

scheduler

Job RID

ic0-b4e7-a0ec3a42f738

Build ID

cd-8748-74e261b07670

Schedule RID

13-8bd5-dac10193cabf

5. Select **Rollback to transaction**.
6. A confirmation dialog will be displayed.

Rollback dataset

Confirm rollback of dataset?

Dataset will revert back to the selected transaction.

Dataset

dataset

Transaction time

Sun, Feb 16, 2025, 3:10 PM

The logic backing the dataset will be left unchanged and will need to be updated to apply to next build.

☒ I acknowledge that rolling back this dataset cannot easily be undone

Rollback dataset

Cancel

7. Acknowledge the warning that a rollback cannot easily be undone and select **Rollback dataset**.

API Reference ↗ is complete, navigate to the dataset's **History** tab and ensure that the actions are now crossed out, as shown below:

Send feedback

← dataset

🔍
🔍
🔍

Preview
SQL scratchpad
History
</> Code
Build timeline
Data health

Feb 16, 3:13 PM
Update • scheduler
Part of build • 1 job

Feb 16, 3:12 PM
Update • scheduler
Part of build • 1 job

Feb 16, 3:11 PM
Update • scheduler
Part of build • 1 job

Feb 16, 3:10 PM
Update • scheduler
Part of build • 1 job

Feb 16, 3:09 PM
Update • scheduler
Part of build • 1 job

Feb 16, 3:08 PM
Update • scheduler
Part of build • 1 job

Feb 16, 3:07 PM
Update • scheduler
Part of build • 1 job

Summary
Overview
Logs
Files
Metadata
Schema
Job spec
Advanced

Transaction details

Status ✔ Committed

Type Update

Duration 55s

Start time Feb 16, 2025, 3:09 PM

Close time Feb 16, 2025, 3:10 PM

Files 2

Size of files 625B

Started by Foundry

Transaction RID fe-b523-2ebbc0aa2bbc

Rollback Rollback to transaction

Job details View schedule View build report

Status ✔ Succeeded (Part of a build)

Type transforms python-1

Duration 53s (part of a 57s build)

Start time Feb 16, 2025, 3:09 PM

End time Feb 16, 2025, 3:10 PM

Started by scheduler

Job RID ic0-b4e7-a0ec3a42f738

Build ID cd-8748-74e261b07670

Schedule RID 13-8bd5-dac10193cabf

⚠ If a dataset backs an [object type](#) stored using Object Storage v2, manual intervention is required to ensure that the object type is [reindexed](#) with a successful run of the [replacement pipeline](#) to reflect the state after the rollback.

Force a snapshot on the next build

i Forcing a snapshot will **not** change the dataset's transaction history or produce immediate visible changes. The snapshot will occur on the next build.


Forcing a snapshot **will** require a force build to rebuild the dataset if there are no changes to either the input data or the logic backing the dataset.

1. [Navigate to a Data Lineage](#) graph containing the dataset you would like to force a

[API Reference](#)

2. Select the dataset in the graph. Then, from the branch selector at the bottom, select the branch on which you would like to perform the rollback.


Send feedback



3. Select the **History** tab in the panel at the bottom of the page. Do not select any transactions.



4. In the displayed **Summary** panel, select the **Force snapshot** option in the toolbar at the top right.



5. A confirmation dialog will be displayed.

[API Reference ↗](#)

[Send feedback](#)

Queue snapshot

Force snapshot on next build

Set the next build (manual or scheduled) of this dataset to snapshot.

Dataset

dataset

The logic backing this dataset will remain unchanged and will need to be updated to apply to the next build.

☐

I acknowledge that queueing a snapshot on this dataset cannot be undone

Queue snapshot

Cancel

6. Acknowledge the warning that a forcing a snapshot on the next build cannot be undone and select **Queue snapshot**.

PREVIOUS

← Roll back a pipeline

NEXT

Understand permissions / Check → resource permissions

© 2026 Palantir Technologies Inc. All rights reserved.

[Cookies Statement ↗](#)

[Privacy Statement ↗](#)

[Terms of Use ↗](#)

—[Cookies Settings](#)


[API Reference ↗](#)

[Send feedback](#)

Data connectivity & integration / Data Lineage / Understand permissions / Check resource permissions

Check resource permissions

You can use Data Lineage to check users' permissions to view datasets or artifacts using the "Permissions" coloring option. To do that, start by adding nodes to the graph. You can do so using the search helper on the side panel.




Then expand the graph to view the lineage leading to your resource (read more about [exploring lineage](#)).

[API Reference ↗](#)




[Send feedback](#)



Once you have done this, use the **Node color options** dropdown to select the **Permissions** color scheme.

[API Reference ↗](#)


[Send feedback](#)



Select the user's name from the **View as** dropdown. This will allow you to see the user's permissions to each of the nodes on the graph.


[API Reference ↗](#)

[Send feedback](#)



There are two permission types you can color by:

- [Data access in datasets](#)
- [Resource access](#)



[API Reference](#) ↗

Data access in datasets

[Send feedback](#)

Use this option to troubleshoot permissions issues. Remember that a user's data access is affected by data lineage (see [Platform Security](#)). By coloring your nodes based on the user's access to data, you can easily see what the upstream datasets are that may restrict the user's access to data.

Note that this option only works on dataset nodes.

Resource access

This will allow you to see the [role](#) (such as Editor, Viewer, etc.) that is set for the selected user on the selected resource.

Use this option to view the level of access users have to your artifacts.

i Roles do not correspond to data lineage the same way that data access does. For example, user being an "Editor" on a Contour Analysis does not guarantee they have permissions to see the data that the analysis depends on. Make sure your users can access the underlying data when sharing a resource with them.

PREVIOUS

← [Understand and manage datasets / Roll back a dataset](#)

NEXT

[See the impact of marking changes](#) →

© 2026 Palantir Technologies Inc. All rights reserved.

[Cookies Statement](#) ↗

[Privacy Statement](#) ↗

[Terms of Use](#) ↗

[API Reference](#) ↗

[Send feedback](#)

Data connectivity & integration / Data Lineage / Understand permissions / See the impact of marking changes

See the impact of Marking changes

You can use Data Lineage to evaluate how changes to dataset Markings can impact derived datasets. This can be useful when [removing Markings](#).

⚠ Warning

Marking simulation relies on the most recent dataset builds and does not account for changes that are not yet finalized. Confirm that you are working with the most up-to-date version of your data.


Access simulation mode

1. Open the **Access information** side panel.
2. Toggle on **Simulate access requirements**.
3. Select any dataset on the graph.
4. Click **Edit markings**.


API Reference ↗



Send feedback



Simulate Marking changes



To simulate Marking application, search for the Marking you want to apply, check the box next to the Marking, and then select the **Simulate changes** button.


Markings that are already applied on a dataset will appear as selected. To simulate Marking removal, uncheck the box next to the Marking and click **Simulate changes**.

API Reference ↗

Remove Markings that were applied directly on the dataset. Removal of Markings that were inherited through a dataset's lineage or from the Data Catalog cannot be simulated.

Send feedback

Analyze the simulated graph



When in simulation mode, the graph coloring will indicate the datasets affected by the Marking changes. The graph colors are labeled in the interface and can represent the following dataset statuses:

- **Simulate changes applied** appears on the datasets to which you applied changes.
- **Access affected** appears on datasets for which the Markings before and after the change will be different.
- **Access unaffected** appears on datasets for which the Markings before and after the change will remain the same.
- **No visible transactions** appears on datasets that have not been built yet, or where you do not have permission to see transactions.

By selecting any of the datasets, the **Access information** side panel will show the simulated access requirements. You can toggle simulation mode on and off to view differences without losing any of the simulation changes.

Tips for understanding changes


Before making changes, we suggest consulting the [Markings documentation](#) to learn more about the impact of Markings on users.

When making Markings, consider the following:

[API Reference ↗](#)

[Send feedback](#)

- Datasets can [stop propagating Markings via code](#).



- In the **Permissions** coloring, nodes on the Data Lineage graph that stop propagating Markings show that data access was *modified via code*. This message will also appear in the **Access information** section of the node properties side panel.
- In the Code Helper, you can check the code for a dataset to see if it stops propagating Markings by using the term `stop_propagating`.
- Datasets can have Markings propagated to them from *other inputs*; expand the dataset inputs by clicking on the left arrow in the dataset node.
- Markings can be applied on the *parent Project or folder*; Markings will have a folder icon on their left when simulation mode is not enabled, and will show a folder icon in the Marking simulation menu when simulation mode is enabled.

PREVIOUS

← [Check resource permissions](#)

NEXT

[Dataset Preview / Overview](#) →

© 2026 Palantir Technologies Inc. All rights reserved.

[Cookies Statement](#) ↗

[Privacy Statement](#) ↗

[Terms of Use](#) ↗

—[Cookies Settings](#)

[API Reference](#) ↗

[Send feedback](#)