


Table of Contents

- 1. Platform overview • Overview • Palantir
- 2. Platform overview • Overview • Palantir
- 3. Platform overview • Architecture • Palantir
- 4. Platform overview • Interoperability • Palantir
- 5. Platform overview • AIP capabilities • Palantir
- 6. Platform overview • Development lifecycle • Palantir



Platform overview


Palantir AIP powers real-time, AI-driven decision-making in the most critical commercial and government contexts around the world. From [public health ↗](#) to [battery production ↗](#), organizations depend on Palantir to safely, securely, and effectively leverage AI in their enterprises — and [drive operational results ↗](#).

In short, Palantir AIP connects generative AI to operations. Together with Foundry - Palantir's data operations platform - and Apollo - Palantir's mission control for autonomous software deployment, AIP is part of an *AI Mesh* that can deliver the full gamut of AI-driven products, from LLM-powered web applications to mobile applications using vision-language models to edge applications that embed localized AI. We call this entire set of capabilities, functionality, and tooling the *Palantir platform*.

API Reference ↗



Send feedback



While many factors contribute to achieving and scaling operational impact with the Palantir platform – including [AIP Bootcamps](#), where customers are hands-on-keyboard and achieving outcomes with AI in a matter of hours – the key differentiator is a software architecture which revolves around the Palantir [Ontology](#).

i Palantir Learning portal

Learn important platform concepts in the ["Introduction to Foundry & AIP for Enterprise Organizations" course on learn.palantir.com](#), or read on below.

The Ontology

The Ontology is designed to represent the *decisions* in an enterprise, not simply the data. Every organization in the world is faced with the challenge of how to execute the best possible decisions, often in real-time, while contending with internal and external conditions that are constantly in flux.

The complexity of these decision processes is reflected in the Ontology, which facilitates deep, two-way interoperability with existing enterprise systems. The Ontology automatically

integrates data, logic and action components into a modern, AI-accessible environment. This unlocks the rapid development of operational applications with AI teaming, in addition to conventional business intelligence and

[API Reference](#)

[Send feedback](#)

Decision components

Every decision can be broken down into **data**, **logic**, and **actions**.

- **Data:** What are the relevant facts or truth about the world and our operations that form the context for this decision?
- **Logic:** What organizational or business rules act as guardrails for this decision? What are the probabilities of certain outcomes under different assumptions? What have we done in previous, similar situations and what have the outcomes been? What are the inputs from our forecasting and optimization models?
- **Actions:** What are the "kinetics" or effects of this decision – that is, how does the decision manifest in the world? How do we reduce or collapse the steps between taking a decision in AIP and affecting an outcome in a production setting?

In the Palantir platform, all of these components are designed to facilitate AI teaming patterns to unlock the full potential of your operators, analysts, and subject-matter experts.

Suggested reading


To learn more about how these decision components interact to guide workflow development, refer to the documentation on [distilling functional requirements](#) as part of the [use case lifecycle](#), or find examples of industry-specific end-to-end workflows in the [AIP Now showcase ↗](#).

Data

[API Reference ↗](#)

[Send feedback](#)

Data Connectivity for Human+AI Teaming



The Ontology integrates data as objects and links in order to make the real-world complexity of operations understandable for both humans and AI. This unlocks the ability to build *Human + AI teaming* workflows.

The Ontology natively supports a wide range of data types as well as a number of extended primitives, such as semantic search for unlocking unstructured data, media references for working with images and video, and value types for embedding additional constraints and context into data. These are the data building blocks for AI workflow development, described further in the Logic and Actions sections below.

This data model powers out-of-the-box applications for exploring structured, unstructured, geospatial, temporal, simulated, and other data modalities. These baseline tools are enriched with the context-aware AIP Assist to dramatically shorten the time-to-value when exploring and analyzing data in the platform.

In [API Reference](#) ↗, you can find information about building and analytics, modeling data in the Ontology automatically creates a robust API gateway and Ontology Software Developer Kit (OSDK) to serve as an “operational bus” for connectivity throughout the enterprise.

[Send feedback](#)

Data connectivity

Data rarely comes packaged in the clean, correct, and well-shaped formats needed to accurately and reliably present truth to decision makers. To that end, the Palantir platform provides an extensible, multimodal data connection and integration framework that works with enterprise data systems out-of-the-box.

[Pipeline Builder](#) puts the power of [LLM data transformation](#) into a point-and-click package, making it easy to use the latest LLMs to power pipeline-based transforms such as classification, sentiment analysis, summarization, entity extraction, or translation. This sets the stage for automatically creating "proposals" in the Ontology for operators to review and approve, without the lag of always running a live request to a model. (Note that, as discussed in the [Logic](#) section below, these two approaches to interacting with models are highly complementary.)

In addition, [AIP Assist](#) in [Pipeline Builder](#) and [Code Repositories](#) accelerates data engineering with an AI partner that not only has access to Palantir documentation and repositories of generic code snippets, but also is deeply integrated into the platform frontend and can suggest next actions or relevant tutorials.

Logic

[API Reference ↗](#)


[Send feedback](#)

Logic Connectivity for Human+AI Teaming

AI + enterprise logic
Anchor AI activity in deterministic, conventional algorithms and business rules

Logic chaining
Ability to chain different forms of logic, within single workflows

"Tool" building framework
Ensure that new logic can be automatically surfaced to AI



If data defines the context for our decisions, logic encapsulates the reasoning and analysis that enriches this context, enabling Human+AI teams to make better decisions. This can be provided as additional context in the form of model outputs and visualizations presented within an operational application, or baked directly into the mechanics of an Action.

Given this broad definition, the ability to define and execute logic shows up throughout the platform; for example, let us consider [models](#), [business logic](#), and [templated analyses and reports](#).

Models

Generative AI, LLMs, forecasts, optimizers, etc.

Models like LLMs or forecasts take parameters and provide an output to serve as context for the decision at hand. In a cycle familiar to data scientists, these models often undergo an iterative training and refinement; however, using these models as operational workloads can be a challenge. Palantir's modeling capabilities can facilitate operational deployment of models.

API Reference ↗

Send feedback

In the Palantir platform, the full lifecycle of a model is captured as a [modeling objective](#) and the logic of the model itself is abstracted with a [model adapter](#). This approach means whether you [train in the platform](#), [bring your own container](#), or [upload a pre-trained model](#), models of all varieties can be bound to the Ontology through [Functions](#) for live interaction embedded in operational apps, or configured for [batch deployment](#) and scheduled for execution in a data pipeline.

Specifically for generative AI, Palantir's [Language Model Service](#) provides a unified interface for multi-modal interactivity while abstracting the specific model and provider implementation details, making it simple to develop across the landscape of commercially-available LLMs. To further improve outcomes, Palantir's [Evaluations](#) tool enables you to benchmark LLM performance over time and between models to monitor drift and make changes with confidence.

Business logic

Business rules, process mapping, semantic search

Where modeling approaches take a "bottom-up" approach to training on data, business logic generally goes "top-down" based on the explicit or implicit rules that govern an operational domain. These may live on an external system, to which Palantir can connect directly with [External Functions](#) and [Webhooks](#) for live interactions in operational workflows, or through [External Transforms](#) for pipeline connectivity. Business logic may also be authored directly within the Palantir platform using [Rules](#) and Pipeline Builder for logic in data pipelines, and Automate and Functions for logic that will be executed live.

Templated analyses and reports


Object views, analysis templates, generated reports

Logic does not live only in data science models or as hard-coded business rules; analysts often capture and collect high-value logic in one-off investigations, analyses, or reports. In the Palantir platform, you can build analyses and dashboards with point-and-click analysis tools like [Contour](#) and [Quiver](#), and notebooks like [Code Workspaces](#). The semantics of the [Ontology data model](#) make it easy to template these analytical products and reuse them, with [API Reference](#) [in Object Views](#) or [Workshop applications](#), or presented as stand-alone [dashboards](#). These object views, templated analyses, and dashboards can be plugged into operational apps to provide at-a-glance insights to guide decision-making and explore new avenues for further ad-hoc exploration.

Send feedback

Taken together, these three facets of logic - models, business logic, and templated analyses and reports - provide a toolkit or palette from which users can mix-and-match to provide decision makers with all of the context needed at the critical moment.

Actions



For any decision to have an impact, the decision must propagate into the world. This is where Actions define the "verbs" of the enterprise - the things that are done - and control how human operators or AI agents can ensure that their decision persists, either within the Ontology data model or through interaction with external systems. In addition, capturing decision outcomes in the Ontology allows users to pair a particular decision with observations of the results in future data. This enables feedback loops that put future decisions in the context of past choices and can be used to retrain or fine-tune models, or simply support operators with a clearer picture of the past.

The atomic unit for representing these "kinetics" in the Ontology is an Action, which provides control for changing or creating data, as well as for orchestrating changes in external systems. Basic actions are simple to define with a point-and-click configuration interface. Actions of arbitrary complexity can be specified

[API Reference](#) ↗

[Send feedback](#)

[backed Actions](#) and the [Ontology Edits Typescript API](#). Actions are also available to package within the [Ontology Software Development Kit](#) (OSDK) and within the [platform API](#) so that custom application development and existing third-party tools can easily and securely write back to the Ontology.

[Permissions](#) for each Action determine which user or agent, under what conditions, is able to execute the action laying the foundation at the lowest level for secure, auditable, and transparent control.

In complex, tightly-coupled environments, such as a supply chain or a manufacturing floor, a small change can cause cascading effects with unexpected or unintended outcomes. The [Scenario](#) primitive allows users to project these consequences by making changes to a branch of the ontology, effectively creating a sandbox universe in which forecasts, business process models, and other analyses can be made on top of (and downstream of) the potential change. The [Vertex](#) application specializes in this kind of process visualization and scenario testing; the [Workshop](#) application builder natively supports scenarios for developing operational applications that incorporate “What if...” workflows.

These primitives create an environment for safe development of Human+AI teams operating in production workflows. The granular permissions and access control from Actions provide a "control plane" in which agents are sandboxed with specific limitations on the data and tools they can wield. In most patterns, rather than directly make changes, AI agents create proposals either synchronously through direct integration with [AIP Logic](#) functions integrated into Workshop, or asynchronously through [Automate](#) or the [Use LLM](#) node in Pipeline Builder. The resulting proposal can then be surfaced to an operator for refinement, feedback, and a resulting decision. This proposal-based pattern, in addition to reinforcing the “human in the loop” paradigm, also generates valuable metadata that enables a positive cycle where the Agent can learn and evolve with continuous feedback.

What's next?

The best way to experience the power of AIP is to start building. Read the [getting started](#) guide for more information, or - if you have access to the platform - just ask AIP Assist where to start based on your intended goals.

[API Reference ↗](#)

 Palantir Learning portal

[Send feedback](#)

To jumpstart building your first end-to-end sample workflow, navigate to learn.palantir.com ↗.

For details on how various platform decision components interact to guide workflow development, refer to the discussion about [distilling functional requirements](#) in the discussion of [use case development](#), or find examples of industry-specific end-to-end workflows in the [AIP Now showcase](#) ↗.

Additionally, you can learn more about how AIP is built and how it integrates with existing investments across your organization:

- [Architecture](#)
- [Interoperability](#)
- [Development lifecycle](#)

Platform capabilities

The remainder of the documentation is organized as a collection of [platform capabilities](#). Summaries of each can be found below:

- [Data connectivity & integration](#)
- [Model connectivity & development](#)
- [Ontology building](#)
- [Use case development](#)
- [Analytics](#)
- [Product delivery](#)
- [Security & governance](#)

Data connectivity & integration

Palantir provides an extensible, multimodal data connection framework that connects to external systems out-of-the-box and provides:

[API Reference](#) ↗

- In-place, zero-copy access to existing data lakes and platforms;

[Send feedback](#)

- An auto-scaling, Kubernetes-based build system for data that works across batch and streaming pipelines;
- Integrated pipeline scheduling and orchestration;
- Native health checks for all data flows; and
- Comprehensive security functionality that spans role-, classification-, and purpose-based access controls.

Suggested reading

Learn more with these resources:

- [Overview: Data connectivity & integration](#)
- [Connecting to data](#)
- [Virtual Tables](#)
- [What is a data pipeline?](#)
- [Health checks](#)
- [Securing a data foundation](#)

Model connectivity & development

Palantir offers an integrated, end-to-end environment for model development (e.g., in Python and R); flexible integration of external models built using industry-standard toolsets; governed paths to production for all developed or integrated models; and a “mission control” for continuous evaluation of deployed models. The architectural goal is to provide a connection path for all business logic and modeling in the enterprise, regardless of where the given asset was trained, tested, and/or hosted.

Suggested reading

Learn more with these resources:

- [Overview: Model connectivity & enablement](#)
- [Model development](#)
- [Integrating external models](#)
- [Evaluating models](#)

API Reference ↗

Send feedback

- [Models in the Ontology](#).

Ontology building

As mentioned above, to create a comprehensive decision-centric model of the enterprise, the Ontology integrates:

- Data, as [objects and links](#);
- Logic, as [models](#) and [functions](#); and
- Actions, as platform [actions](#).

These building blocks of the Ontology make the real-world complexity of operations understandable to both operators and AI, unlocking the ability to build hybrid human-AI workflows. Additional capabilities include:

- Structured mechanisms for capturing data from end users back into the semantic foundation;
- Out-of-the-box applications for exploring the Ontology in structured, unstructured, geospatial, temporal, simulated, and other paradigms; and
- The Ontology Software Developer Kit (OSDK) for leveraging the Ontology as an “operational bus” throughout all parts of the enterprise.

Suggested reading

Learn more with these resources:

- [Overview: Ontology building](#)
- [Why create an Ontology?](#)
- [Ontology and AIP Logic](#)
- [Models in the Ontology](#).
- [Ontology-aware applications](#)
- [System graphs and simulations](#)

[Ontology SDK](#)

[API Reference](#) ↗

Send feedback

Use case development

Palantir's application development framework enables enterprises to build operational workflows and develop use cases that leverage user actions, alerting, and other end-user frontline functions in collaboration with tool-wielding, data-aware AIP Chatbots.

Use case development capabilities include:

- Integration with AIP Logic for building custom workflow agents;
- AI-assisted, low-code / no-code application building that automates security enforcement and the management of underlying storage and compute as well as data and model bindings;
- An application development framework with live preview; and
- APIs, webhooks, and other interfaces that allow for full-spectrum integration with the enterprise.

Suggested reading

Learn more with these resources:

- [Overview: Use case development](#)
- [Creating operational applications](#)
- [Connecting analytics to operations](#)
- [Low-code and no-code application building](#)
- [Integrating with existing systems](#)

Analytics

The platform provides analytical capabilities for every type of user, whether they can code or not. Capabilities include both point-and-click and code-based tools that enable table-based analysis, top-down visual analysis, geospatial analysis, time series analysis, scenario simulation, and more.

[API Reference ↗](#) Palantir's platform goes beyond conventional “read-only” paradigms to write data back into the Ontology, producing valuable new insights within unified governance models.

[Send feedback](#)

The platform also interoperates with common modeling environments (supporting native usage of JupyterLab® and RStudio® Workbench with [Code Workspaces](#)) and business intelligence platforms (including [dedicated connectors](#) for Tableau® and PowerBI®).

Suggested reading

Learn more with these resources:

- [Overview: Analytics](#)
- [Types of analysis](#)
- [Code-based analysis](#)
- [Time series analysis](#)

Product delivery

The Palantir platform provides DevOps tooling to package, deploy, and maintain data products built in the platform. These product delivery capabilities include a packaging interface to create "products" consisting of collections of platform resources (pipelines, Ontologies, applications, models, etc.); a Marketplace storefront for product discovery and installation; and the ability to manage product installations with automatic upgrades, maintenance windows, and more.

Suggested reading

Learn more with these resources:

- [Overview: Product delivery](#)
- [Core concepts](#)
- [Finding and installing products in Marketplace](#)
- [Creating products](#)

Security and Governance

[API Reference](#) ↗

The Palantir platform features a comprehensive, best-in-class security model that propagates across the entire platform and, by default, remains with information wherever it

[Send feedback](#)

travels. Capabilities include:

- [Encryption of all data](#), both in transit and at rest;
- Authentication and identity protection controls;
- Authorization controls that can blend role-, marking-, and purpose-driven paradigms;
- Robust security [audit logging](#); and
- Highly extensible information governance, management, and [privacy controls](#).



Suggested reading

Learn more with these resources:

- [Overview: Security & governance](#)
- [Enterprise security](#)
- [Audit logging](#)
- [Privacy controls](#)

Management & enablement

Platform administrators have access to a robust set of tools for managing the Palantir platform. The core applications for platform management are:

- [Control Panel](#)
- [Resource Management](#)
- [Upgrade Assistant](#)

Platform administrators and program managers also have access to resources to facilitate user enablement, such as [AIP Assist](#). These resources are described in the [management & enablement documentation](#).



Suggested reading

Learn more with these resources:

API Reference ↗

- [Overview: Management & enablement](#)
- [Enrollments and organizations](#)



Send feedback

- [Authentication](#)
- [Retention](#)

← PREVIOUS
Getting help / HTTP errors

NEXT →
Architecture

© 2026 Palantir Technologies Inc. All rights reserved.

[Cookies Statement ↗](#)


[Privacy Statement ↗](#)

[Terms of Use ↗](#)

— [Cookies Settings](#)

[API Reference ↗](#)

[Send feedback](#)



Platform overview


Palantir AIP powers real-time, AI-driven decision-making in the most critical commercial and government contexts around the world. From [public health ↗](#) to [battery production ↗](#), organizations depend on Palantir to safely, securely, and effectively leverage AI in their enterprises — and [drive operational results ↗](#).

In short, Palantir AIP connects generative AI to operations. Together with Foundry - Palantir's data operations platform - and Apollo - Palantir's mission control for autonomous software deployment, AIP is part of an *AI Mesh* that can deliver the full gamut of AI-driven products, from LLM-powered web applications to mobile applications using vision-language models to edge applications that embed localized AI. We call this entire set of capabilities, functionality, and tooling the *Palantir platform*.

API Reference ↗



Send feedback



While many factors contribute to achieving and scaling operational impact with the Palantir platform — including [AIP Bootcamps](#), where customers are hands-on-keyboard and achieving outcomes with AI in a matter of hours — the key differentiator is a software architecture which revolves around the Palantir [Ontology](#).

i Palantir Learning portal

Learn important platform concepts in the ["Introduction to Foundry & AIP for Enterprise Organizations" course on learn.palantir.com](#), or read on below.

The Ontology

The Ontology is designed to represent the *decisions* in an enterprise, not simply the data. Every organization in the world is faced with the challenge of how to execute the best possible decisions, often in real-time, while contending with internal and external conditions that are constantly in flux.

The complexity of these decision processes is reflected in the Ontology, which facilitates deep, two-way interoperability with existing enterprise systems. The Ontology automatically

integrates data, logic and action components into a modern, AI-accessible environment. This unlocks the rapid development of operational applications with AI teaming, in addition to conventional business intelligence and

[API Reference](#)

[Send feedback](#)

Decision components

Every decision can be broken down into **data**, **logic**, and **actions**.

- **Data:** What are the relevant facts or truth about the world and our operations that form the context for this decision?
- **Logic:** What organizational or business rules act as guardrails for this decision? What are the probabilities of certain outcomes under different assumptions? What have we done in previous, similar situations and what have the outcomes been? What are the inputs from our forecasting and optimization models?
- **Actions:** What are the "kinetics" or effects of this decision – that is, how does the decision manifest in the world? How do we reduce or collapse the steps between taking a decision in AIP and affecting an outcome in a production setting?

In the Palantir platform, all of these components are designed to facilitate AI teaming patterns to unlock the full potential of your operators, analysts, and subject-matter experts.

Suggested reading


To learn more about how these decision components interact to guide workflow development, refer to the documentation on [distilling functional requirements](#) as part of the [use case lifecycle](#), or find examples of industry-specific end-to-end workflows in the [AIP Now showcase ↗](#).

Data

[API Reference ↗](#)

[Send feedback](#)

Data Connectivity for Human+AI Teaming



The Ontology integrates data as objects and links in order to make the real-world complexity of operations understandable for both humans and AI. This unlocks the ability to build *Human + AI teaming* workflows.

The Ontology natively supports a wide range of data types as well as a number of extended primitives, such as semantic search for unlocking unstructured data, media references for working with images and video, and value types for embedding additional constraints and context into data. These are the data building blocks for AI workflow development, described further in the Logic and Actions sections below.

This data model powers out-of-the-box applications for exploring structured, unstructured, geospatial, temporal, simulated, and other data modalities. These baseline tools are enriched with the context-aware AIP Assist to dramatically shorten the time-to-value when exploring and analyzing data in the platform.

In [API Reference](#) ↗, you can find information about building and analytics, modeling data in the Ontology automatically creates a robust API gateway and Ontology Software Developer Kit (OSDK) to serve as an “operational bus” for connectivity throughout the enterprise.

[Send feedback](#)

Data connectivity

Data rarely comes packaged in the clean, correct, and well-shaped formats needed to accurately and reliably present truth to decision makers. To that end, the Palantir platform provides an extensible, multimodal data connection and integration framework that works with enterprise data systems out-of-the-box.

[Pipeline Builder](#) puts the power of [LLM data transformation](#) into a point-and-click package, making it easy to use the latest LLMs to power pipeline-based transforms such as classification, sentiment analysis, summarization, entity extraction, or translation. This sets the stage for automatically creating "proposals" in the Ontology for operators to review and approve, without the lag of always running a live request to a model. (Note that, as discussed in the [Logic](#) section below, these two approaches to interacting with models are highly complementary.)

In addition, [AIP Assist](#) in [Pipeline Builder](#) and [Code Repositories](#) accelerates data engineering with an AI partner that not only has access to Palantir documentation and repositories of generic code snippets, but also is deeply integrated into the platform frontend and can suggest next actions or relevant tutorials.

Logic

[API Reference ↗](#)


[Send feedback](#)

Logic Connectivity for Human+AI Teaming

AI + enterprise logic
Anchor AI activity in deterministic, conventional algorithms and business rules

Logic chaining
Ability to chain different forms of logic, within single workflows

"Tool" building framework
Ensure that new logic can be automatically surfaced to AI



If data defines the context for our decisions, logic encapsulates the reasoning and analysis that enriches this context, enabling Human+AI teams to make better decisions. This can be provided as additional context in the form of model outputs and visualizations presented within an operational application, or baked directly into the mechanics of an Action.

Given this broad definition, the ability to define and execute logic shows up throughout the platform; for example, let us consider [models](#), [business logic](#), and [templated analyses and reports](#).

Models

Generative AI, LLMs, forecasts, optimizers, etc.

Models like LLMs or forecasts take parameters and provide an output to serve as context for the decision at hand. In a cycle familiar to data scientists, these models often undergo an iterative training and refinement; however, using these models as operational workloads can be a challenge. Palantir's modeling capabilities can facilitate operational deployment of models.

API Reference ↗

Send feedback

In the Palantir platform, the full lifecycle of a model is captured as a [modeling objective](#) and the logic of the model itself is abstracted with a [model adapter](#). This approach means whether you [train in the platform](#), [bring your own container](#), or [upload a pre-trained model](#), models of all varieties can be bound to the Ontology through [Functions](#) for live interaction embedded in operational apps, or configured for [batch deployment](#) and scheduled for execution in a data pipeline.

Specifically for generative AI, Palantir's [Language Model Service](#) provides a unified interface for multi-modal interactivity while abstracting the specific model and provider implementation details, making it simple to develop across the landscape of commercially-available LLMs. To further improve outcomes, Palantir's [Evaluations](#) tool enables you to benchmark LLM performance over time and between models to monitor drift and make changes with confidence.

Business logic

Business rules, process mapping, semantic search

Where modeling approaches take a "bottom-up" approach to training on data, business logic generally goes "top-down" based on the explicit or implicit rules that govern an operational domain. These may live on an external system, to which Palantir can connect directly with [External Functions](#) and [Webhooks](#) for live interactions in operational workflows, or through [External Transforms](#) for pipeline connectivity. Business logic may also be authored directly within the Palantir platform using [Rules](#) and Pipeline Builder for logic in data pipelines, and Automate and Functions for logic that will be executed live.

Templated analyses and reports


Object views, analysis templates, generated reports

Logic does not live only in data science models or as hard-coded business rules; analysts often capture and collect high-value logic in one-off investigations, analyses, or reports. In the Palantir platform, you can build analyses and dashboards with point-and-click analysis tools like [Contour](#) and [Quiver](#), and notebooks like [Code Workspaces](#). The semantics of the [Ontology data model](#) make it easy to template these analytical products and reuse them, with [API Reference](#) in [Object Views](#) or [Workshop applications](#), or presented as stand-alone [dashboards](#). These object views, templated analyses, and dashboards can be plugged into operational apps to provide at-a-glance insights to guide decision-making and explore new avenues for further ad-hoc exploration.

Send feedback

Taken together, these three facets of logic - models, business logic, and templated analyses and reports - provide a toolkit or palette from which users can mix-and-match to provide decision makers with all of the context needed at the critical moment.

Actions



For any decision to have an impact, the decision must propagate into the world. This is where Actions define the "verbs" of the enterprise - the things that are done - and control how human operators or AI agents can ensure that their decision persists, either within the Ontology data model or through interaction with external systems. In addition, capturing decision outcomes in the Ontology allows users to pair a particular decision with observations of the results in future data. This enables feedback loops that put future decisions in the context of past choices and can be used to retrain or fine-tune models, or simply support operators with a clearer picture of the past.

The atomic unit for representing these "kinetics" in the Ontology is an Action, which provides control for changing or creating data, as well as for orchestrating changes in external systems. Basic actions are simple to define with a point-and-click configuration interface. Actions of arbitrary complexity can be specified

[API Reference ↗](#)

[Send feedback](#)

[backed Actions](#) and the [Ontology Edits Typescript API](#). Actions are also available to package within the [Ontology Software Development Kit](#) (OSDK) and within the [platform API](#) so that custom application development and existing third-party tools can easily and securely write back to the Ontology.

[Permissions](#) for each Action determine which user or agent, under what conditions, is able to execute the action laying the foundation at the lowest level for secure, auditable, and transparent control.

In complex, tightly-coupled environments, such as a supply chain or a manufacturing floor, a small change can cause cascading effects with unexpected or unintended outcomes. The [Scenario](#) primitive allows users to project these consequences by making changes to a branch of the ontology, effectively creating a sandbox universe in which forecasts, business process models, and other analyses can be made on top of (and downstream of) the potential change. The [Vertex](#) application specializes in this kind of process visualization and scenario testing; the [Workshop](#) application builder natively supports scenarios for developing operational applications that incorporate “What if...” workflows.

These primitives create an environment for safe development of Human+AI teams operating in production workflows. The granular permissions and access control from Actions provide a "control plane" in which agents are sandboxed with specific limitations on the data and tools they can wield. In most patterns, rather than directly make changes, AI agents create proposals either synchronously through direct integration with [AIP Logic](#) functions integrated into Workshop, or asynchronously through [Automate](#) or the [Use LLM](#) node in Pipeline Builder. The resulting proposal can then be surfaced to an operator for refinement, feedback, and a resulting decision. This proposal-based pattern, in addition to reinforcing the “human in the loop” paradigm, also generates valuable metadata that enables a positive cycle where the Agent can learn and evolve with continuous feedback.

What's next?

The best way to experience the power of AIP is to start building. Read the [getting started](#) guide for more information, or - if you have access to the platform - just ask AIP Assist where to start based on your intended goals.

[API Reference ↗](#)

 Palantir Learning portal

[Send feedback](#)

To jumpstart building your first end-to-end sample workflow, navigate to learn.palantir.com ↗.

For details on how various platform decision components interact to guide workflow development, refer to the discussion about [distilling functional requirements](#) in the discussion of [use case development](#), or find examples of industry-specific end-to-end workflows in the [AIP Now showcase](#) ↗.

Additionally, you can learn more about how AIP is built and how it integrates with existing investments across your organization:

- [Architecture](#)
- [Interoperability](#)
- [Development lifecycle](#)

Platform capabilities

The remainder of the documentation is organized as a collection of [platform capabilities](#). Summaries of each can be found below:

- [Data connectivity & integration](#)
- [Model connectivity & development](#)
- [Ontology building](#)
- [Use case development](#)
- [Analytics](#)
- [Product delivery](#)
- [Security & governance](#)

Data connectivity & integration

Palantir provides an extensible, multimodal data connection framework that connects to external systems out-of-the-box and provides:

[API Reference](#) ↗

- In-place, zero-copy access to existing data lakes and platforms;

[Send feedback](#)

- An auto-scaling, Kubernetes-based build system for data that works across batch and streaming pipelines;
- Integrated pipeline scheduling and orchestration;
- Native health checks for all data flows; and
- Comprehensive security functionality that spans role-, classification-, and purpose-based access controls.

Suggested reading

Learn more with these resources:

- [Overview: Data connectivity & integration](#)
- [Connecting to data](#)
- [Virtual Tables](#)
- [What is a data pipeline?](#)
- [Health checks](#)
- [Securing a data foundation](#)

Model connectivity & development

Palantir offers an integrated, end-to-end environment for model development (e.g., in Python and R); flexible integration of external models built using industry-standard toolsets; governed paths to production for all developed or integrated models; and a “mission control” for continuous evaluation of deployed models. The architectural goal is to provide a connection path for all business logic and modeling in the enterprise, regardless of where the given asset was trained, tested, and/or hosted.

Suggested reading

Learn more with these resources:

- [Overview: Model connectivity & enablement](#)
- [Model development](#)
- [API Reference ↗ external models](#)
- [Evaluating models](#)

API Reference ↗

Send feedback

- [Models in the Ontology](#).

Ontology building

As mentioned above, to create a comprehensive decision-centric model of the enterprise, the Ontology integrates:

- Data, as [objects and links](#);
- Logic, as [models](#) and [functions](#); and
- Actions, as platform [actions](#).

These building blocks of the Ontology make the real-world complexity of operations understandable to both operators and AI, unlocking the ability to build hybrid human-AI workflows. Additional capabilities include:

- Structured mechanisms for capturing data from end users back into the semantic foundation;
- Out-of-the-box applications for exploring the Ontology in structured, unstructured, geospatial, temporal, simulated, and other paradigms; and
- The Ontology Software Developer Kit (OSDK) for leveraging the Ontology as an “operational bus” throughout all parts of the enterprise.

Suggested reading

Learn more with these resources:

- [Overview: Ontology building](#)
- [Why create an Ontology?](#)
- [Ontology and AIP Logic](#)
- [Models in the Ontology](#).
- [Ontology-aware applications](#)
- [System graphs and simulations](#)

[Ontology SDK](#)

[API Reference](#) ↗

Send feedback

Use case development

Palantir's application development framework enables enterprises to build operational workflows and develop use cases that leverage user actions, alerting, and other end-user frontline functions in collaboration with tool-wielding, data-aware AIP Chatbots.

Use case development capabilities include:

- Integration with AIP Logic for building custom workflow agents;
- AI-assisted, low-code / no-code application building that automates security enforcement and the management of underlying storage and compute as well as data and model bindings;
- An application development framework with live preview; and
- APIs, webhooks, and other interfaces that allow for full-spectrum integration with the enterprise.

Suggested reading

Learn more with these resources:

- [Overview: Use case development](#)
- [Creating operational applications](#)
- [Connecting analytics to operations](#)
- [Low-code and no-code application building](#)
- [Integrating with existing systems](#)

Analytics

The platform provides analytical capabilities for every type of user, whether they can code or not. Capabilities include both point-and-click and code-based tools that enable table-based analysis, top-down visual analysis, geospatial analysis, time series analysis, scenario simulation, and more.

[API Reference ↗](#) Palantir's platform goes beyond conventional “read-only” paradigms to write data back into the Ontology, producing valuable new insights within unified data governance models.

[Send feedback](#)

The platform also interoperates with common modeling environments (supporting native usage of JupyterLab® and RStudio® Workbench with [Code Workspaces](#)) and business intelligence platforms (including [dedicated connectors](#) for Tableau® and PowerBI®).

Suggested reading

Learn more with these resources:

- [Overview: Analytics](#)
- [Types of analysis](#)
- [Code-based analysis](#)
- [Time series analysis](#)

Product delivery

The Palantir platform provides DevOps tooling to package, deploy, and maintain data products built in the platform. These product delivery capabilities include a packaging interface to create "products" consisting of collections of platform resources (pipelines, Ontologies, applications, models, etc.); a Marketplace storefront for product discovery and installation; and the ability to manage product installations with automatic upgrades, maintenance windows, and more.

Suggested reading

Learn more with these resources:

- [Overview: Product delivery](#)
- [Core concepts](#)
- [Finding and installing products in Marketplace](#)
- [Creating products](#)

 [API Reference](#)  [Governance](#)

The Palantir platform features a comprehensive, best-in-class security model that propagates across the entire platform and, by default, remains with information wherever it

[Send feedback](#)

travels. Capabilities include:

- [Encryption of all data](#), both in transit and at rest;
- Authentication and identity protection controls;
- Authorization controls that can blend role-, marking-, and purpose-driven paradigms;
- Robust security [audit logging](#); and
- Highly extensible information governance, management, and [privacy controls](#).



Suggested reading

Learn more with these resources:

- [Overview: Security & governance](#)
- [Enterprise security](#)
- [Audit logging](#)
- [Privacy controls](#)

Management & enablement

Platform administrators have access to a robust set of tools for managing the Palantir platform. The core applications for platform management are:

- [Control Panel](#)
- [Resource Management](#)
- [Upgrade Assistant](#)

Platform administrators and program managers also have access to resources to facilitate user enablement, such as [AIP Assist](#). These resources are described in the [management & enablement documentation](#).



Suggested reading

Learn more with these resources:

API Reference ↗

- [Overview: Management & enablement](#)
- [Enrollments and organizations](#)



Send feedback

- [Authentication](#)
- [Retention](#)

← PREVIOUS
Getting help / HTTP errors

NEXT →
Architecture

© 2026 Palantir Technologies Inc. All rights reserved.

[Cookies Statement ↗](#)

[Privacy Statement ↗](#)

[Terms of Use ↗](#)

— [Cookies Settings](#)

[API Reference ↗](#)

[Send feedback](#)

Getting started / Platform overview / Architecture

Architecture

Learn more about the architecture of Palantir's platforms in the [Architecture Center](#).

PREVIOUS

← Overview

NEXT

Interoperability →

© 2026 Palantir Technologies Inc. All rights reserved.

[Cookies Statement ↗](#)

[Privacy Statement ↗](#)

[Terms of Use ↗](#)

— [Cookies Settings](#)

[API Reference ↗](#)



Send feedback

Getting started / Platform overview / Interoperability

Interoperability

Learn more about interoperability in the [Architecture Center](#).

PREVIOUS

← [Architecture](#)

NEXT

[AIP capabilities](#) →

© 2026 Palantir Technologies Inc. All rights reserved.

[Cookies Statement ↗](#)

[Privacy Statement ↗](#)

[Terms of Use ↗](#)

— [Cookies Settings](#)

[API Reference ↗](#)



Send feedback

Getting started / Platform overview / AIP capabilities

AIP capabilities

Palantir AIP (Artificial Intelligence Platform) provides a full set of capabilities for building with generative AI and connecting AI to operations. This includes governed access to a wide range of LLMs and multimodal models, building LLM-driven functions, creating and managing agents, managing the evaluation suites that measure AI performance, defining the automations that will run across enterprise processes, and more.

AIP streamlines the deployment, management, and scalability of LLMs while ensuring robust data governance and security. Beyond standalone application builder experiences, AIP transforms the [development toolchain](#), enabling LLMs and other forms of generative AI to power ontology-driven functions and agents.

Learn more about AIP in the [AIP documentation](#).

Note: AIP feature availability is subject to change and may differ between customers.

PREVIOUS
← [Interoperability](#)

NEXT
[Development lifecycle](#) →

© 2026 Palantir Technologies Inc. All rights reserved.

[Cookies Statement](#) ↗

[Privacy Statement](#) ↗

[API Reference](#) ↗



Send feedback

Development lifecycle

At Palantir, we strive to continuously evolve production-grade software to meet the needs of our customers and improve the capabilities we offer. We incorporate cutting-edge technologies and frameworks while remaining grounded in the real-world problems our customers face every day. In practice, this means that Palantir can develop new capabilities in close partnership with customers.

This page outlines the various development life cycle phases of capabilities in the Palantir platform, including the definitions of these statuses, their availability, support levels, and the expectations we set for those statuses between Palantir and our customers.

 While this page explains the typical execution of the product development life cycle at Palantir, some attributes of a life cycle phase may vary based on a given product or feature.

Phases of development: Summary

This table provides an overview of the different development statuses and their core attributes. For more information, see the linked detail sections on each status.

Status	Development	Availability	User support	Public documentation
Beta: Early-stage API Reference ↗ API Reference ↗	Active	Opt-in availability	Yes	Yes Send feedback

Status	Development	Availability	User support	Public documentation
<u>Generally available (GA)</u> : Production feature	● Active	● Broadly available	● Yes	● Yes
<u>Legacy</u> : Production feature without active development	● Critical fixes only	● Broadly available *	● Yes	● Yes
<u>Sunset</u> : Feature slated for deprecation, but no date is scheduled	● Critical fixes only	● Existing customers only	● Yes	● Yes
<u>Planned deprecation</u> : Deprecation date and migration path communicated	● Critical fixes only	● Existing customers only	● Yes	● Yes
<u>Deprecated</u> : Feature is no longer available	● None	● None	● None	● None

* Installed but disabled for new customers by default. New customers can enable the feature through Control Panel.

Phases of development

Beta

Once a feature's purpose and approach has been validated with one or more customer partners, it is made available to a broader group of customers in a **beta** phase.

[API Reference ↗](#)
 phase are to invest heavily in the long-term maintainability of a feature while measuring overall value and usage. For this reason, beta features may be rolled out slowly to customers to gather feedback during the development process. [Send feedback](#)

features are available to a wide range of customers; at other times, beta features may be

limited to partners with specific business needs or feedback opportunities that advance the development process.

Beta features are documented publicly to help enable customer partners to use this early-stage functionality, even if the features are not yet available in every enrollment. When a feature is in the beta phase, it will be labeled as such in documentation.

There is no guarantee that beta features or products will progress to the next stages of development, though the vast majority do move on to general availability.

Generally available (GA)

Generally available is the term used to describe the vast majority of features and products in Foundry. GA features are enabled for customers by default and constitute a core part of the platform. When features become GA, other parts of the platform may build on top of them to enable tight-knit integrations.

In the GA phase, there is a continuous feedback loop between Palantir engineers and the broad set of customers using the feature. Feedback is triaged, prioritized, and fed into Foundry's product roadmap, and updates are delivered to customers rapidly using [Apollo 7](#), Palantir's platform for continuous delivery.

Although GA features are widely available, it is possible in some cases that your enrollment may not have every feature enabled as some features have dependencies on specific types of infrastructure or may require specific contractual agreements. For example, some GA features in Foundry are only available in Palantir's managed SaaS environment and are not supported for self-hosted installations. Keep in mind that when a new feature or application is [announced](#) as generally available, there may be a delay of a week or more before it becomes available within a specific Foundry enrollment.

Unless otherwise specified, any publicly documented feature is generally available, and you can rely on GA features to be fully supported. The removal of any GA feature from the platform will follow the sunset and planned deprecation processes outlined below.

 [API Reference ↗](#)

[Send feedback](#)

In the **legacy** phase of development, features and products enter a stage where work is considered complete and no additional feature development is expected.

Legacy applications are widely installed across the Palantir customer base and available for continued production use. As in GA, products in the legacy status remain fully supported. For new customers, they are typically available on an opt-in basis. Despite legacy features not being actively developed, feedback is still encouraged as it can inform our work on new applications and tools.

All products and features in the legacy phase will be labeled as such and thoroughly documented publicly. At this point, no deprecation date is planned or expected.

Sunset

As development progresses, an existing feature or application in Foundry may reach the end of its usefulness or purpose, or it may be superseded by another functionality. The original vision for the feature may be narrower than the problem space it has grown to address, or new tools or features may result in more robust or scalable solutions to a problem. At this point, we start to discourage use of the feature, and new customers will not be able to enable it for use in their enrollment.

In our public documentation, we will label products that are in the sunset phase and will suggest other tools and applications to replace them. As in GA, products in the sunset phase remain fully supported. Deprecation is planned for the future, but no exact date is confirmed. As soon a deprecation date is identified, the feature will move into the **planned deprecation** phase, with the date and any relevant migration details communicated.

Planned deprecation

Products entering the **planned deprecation** phase are preparing for retirement from the platform. As products or features enter the planned deprecation phase, Palantir works to provide proactive communication around the scheduled retirement date and any migration steps, if needed. Products in planned deprecation remain fully supported. However, new e

[API Reference ↗](#) be able to enable features in planned deprecation.

In planned deprecation, a feature *does* have a confirmed deprecation date and is in the process of removing the feature and its integrations from our platform. Like products in

[Send feedback](#)

the sunset phase, we will label products in the planned deprecation in our public documentation and in the platform. Our public documentation will also offer suggestions for workflow migrations that will fit your use case.

When we move a feature into planned deprecation, you can expect time to migrate your workflows to new tools and applications before the scheduled deprecation date. We will use our platform [Upgrade Assistant](#) to notify administrators of the upcoming deprecation and provide a clear deadline for complying with workflow migrations. For application-level deprecations, the intent to deprecate and the final deprecation notice will be proactively shared with the registered platform administrator's [contact details](#) as well as publicly on the [Foundry Announcements page](#). Progress towards feature deprecation is tracked quantitatively to ensure all customers are able to migrate to a replacement before a feature is finally removed from the platform.

Deprecated

When a product completes the required tasks during the planned deprecation phase, it is considered fully *deprecated*. In this phase, the feature or product is no longer supported or available in the platform and all customers have been migrated away from relevant workflows. Customers will no longer have access to related API endpoints or other integrations, and support is no longer available. All public documentation of the feature is removed and no longer discoverable.

